

# Learning Layout and Style Reconfigurable GANs for Controllable Image Synthesis

Wei Sun and Tianfu Wu

**Abstract**—With the remarkable recent progress on learning deep generative models, it becomes increasingly interesting to develop models for *controllable* image synthesis from *reconfigurable* inputs. This paper focuses on a recent emerged task, *layout-to-image*, to learn generative models that are capable of synthesizing photo-realistic images from spatial layout (*i.e.*, object bounding boxes configured in an image lattice) and style (*i.e.*, structural and appearance variations encoded by latent vectors). This paper first proposes an intuitive paradigm for the task, *layout-to-mask-to-image*, to learn to unfold object masks of given bounding boxes in an input layout to bridge the gap between the input layout and synthesized images. Then, this paper presents a method built on Generative Adversarial Networks for the proposed layout-to-mask-to-image with style control at both image and mask levels. Object masks are learned from the input layout and iteratively refined along stages in the generator network. Style control at the image level is the same as in vanilla GANs, while style control at the object mask level is realized by a proposed novel feature normalization scheme, *Instance-Sensitive and Layout-Aware Normalization*. In experiments, the proposed method is tested in the COCO-Stuff dataset and the Visual Genome dataset with state-of-the-art performance obtained.

**Index Terms**—Image Synthesis; Layout-to-Image; Layout-to-Mask-to-Image; Deep Generative Learning; GAN; ISLA-Norm.

## 1 INTRODUCTION

### 1.1 Motivation and Objective

REMARKABLE recent progress has been made on both unconditional and conditional image synthesis [1], [2], [3], [4], [5], [6], [7], [8]. The former aims to generate high-fidelity images from some random latent codes/vectors (*e.g.*, sampled from the standard multivariate Gaussian distribution). The latter needs to do so with given conditions satisfied in terms of some consistency metrics. The conditions may take many forms such as category labels [3], [9], paired or unpaired source images [10], [11], [12], [13], semantic maps [14], [15], text description [16], [17] and scene graphs [18], [19]. Conditional image synthesis, especially with coarse yet complicated and reconfigurable conditions, remains a long-standing problem. As illustrated in Fig. 1, we shall only focus on conditional image synthesis from spatial layout and style latent codes, so-called *layout-to-image* [20]. Powerful systems, once developed, can pave a way for computers to truly understand visual patterns and their compositions via a comprehensive and systematic “analysis-by-synthesis” scheme. Those systems will also enable a wide range of practical applications, *e.g.*, generating high-fidelity data for long-tail scenarios in different vision tasks such as autonomous driving.

In layout-to-image, the layout that a synthesized image needs to satisfy consists of labeled bounding boxes configured in an image lattice (*e.g.*, 256 × 256 pixels). The style of a synthesized image refers to structural and appearance variations at both image and object levels, which is often encoded by some latent codes. Generating images from a spatial layout represents a sweet spot in conditional image



Fig. 1. Illustration of the *conditional image synthesis* task studied in this paper. Unlike *unconditional image synthesis* which learns to generate images from a latent code and mainly cares about the overall realness of synthesized images, we are interested in learning to synthesize images from an input layout (object bounding boxes configured in an image lattice such as 256 × 256 pixels) with style controlled by input latent codes. In addition to realness, synthesized images are required to satisfy the given layout with multiple plausible realizations. Furthermore, both the layout and style are reconfigurable and a generator model is required to be consistent with the reconfigurations, thus leading to controllable conditional image synthesis. The right shows three examples generated by our proposed method at a resolution of 256 × 256. See text for details.

synthesis. Spatial layouts are usually used as intermediate representations for other conditional image synthesis tasks such as text-to-image [17], [21] and scene-graph-to-image [18], [19]. And, layouts are more flexible, less constrained and easier to collect than other conditions such as semantic segmentation maps [11], [14]. Existing object detection benchmarks can be exploited in training.

The generative learning task of layout-to-image was recently proposed and only a few work have been proposed in the very recent literature [18], [19], [20], [23]. Although relatively new, it has been well recognized in the computer vision community. For example, the work (Grid2Im) by Ashua and Wolf [19] won the best paper honorable mentions at ICCV 2019. The layout-to-image task was emerged under the context of remarkable progress made in conditional image synthesis with relatively less complicated conditions such as the class-conditional image synthesis in ImageNet by the BigGAN [5], and the amazing style control for specific

• W. Sun and T. Wu are with the Department of Electrical and Computer Engineering and the Visual Narrative Initiative, North Carolina State University, USA.  
E-mail: {wsun12, tianfu\_wu}@ncsu.edu

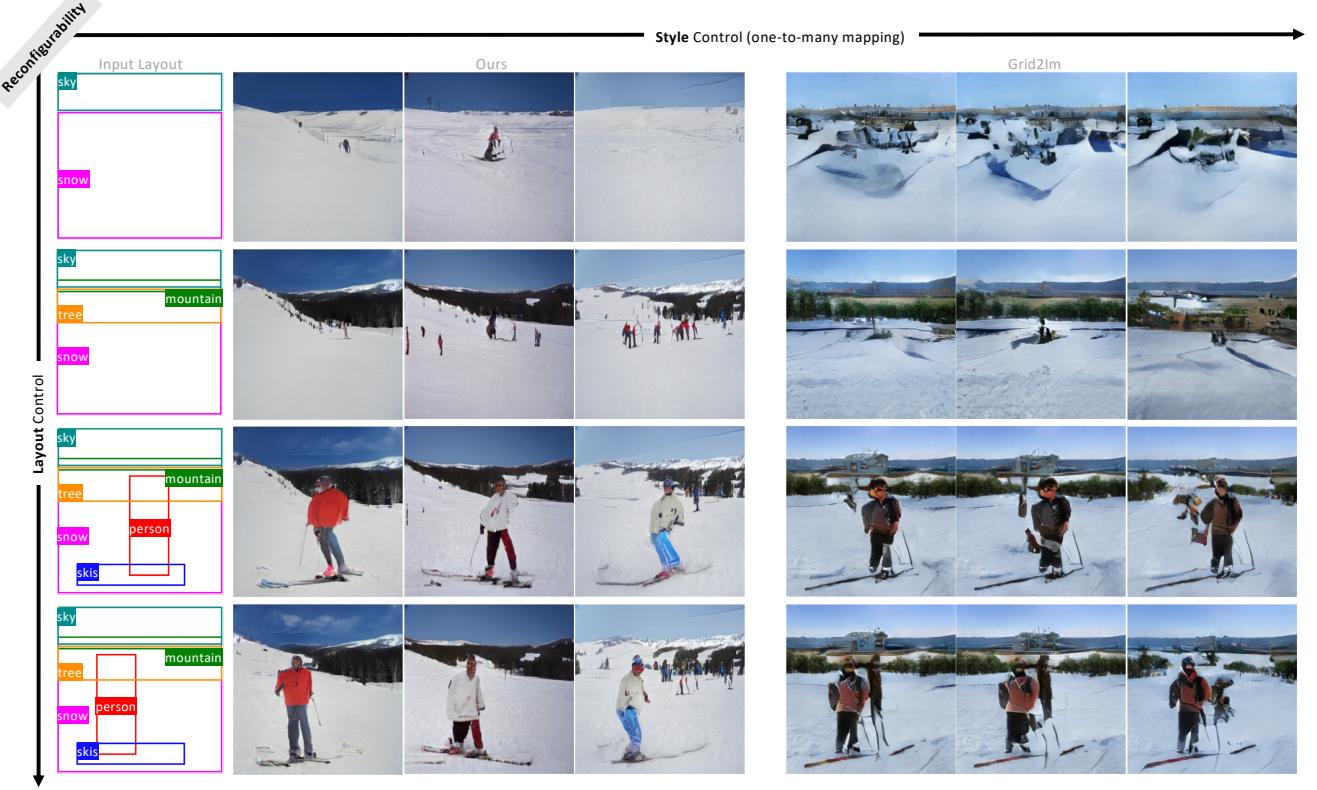


Fig. 2. Illustration of controllable image synthesis from reconfigurable spatial layouts and style codes. The proposed method is compared with the prior art, the Grid2Im by Ashua and Wolf [19]. *Each row shows effects of style control*, in which three synthesized images are shown using the same input layout on the left by randomly sampling three style latent codes. *Each column shows effects of layout control* in terms of consecutively adding new objects (the first three) or perturbing an object bounding box (the last one), while retaining the style codes of existing objects unchanged. Compared to Grid2Im, the proposed method can generate more diverse images with respect to style control (e.g., the appearance of snow, and the pose and appearance of person). The proposed method also shows stronger controllability in retaining the style between consecutive spatial layouts. For example, in the second row, the snow region is not significantly affected by the newly added mountain and tree regions. Our method can retain the style of snow very similar, while the Grid2Im seems to fail to control. Similarly, between the last two rows, our method can produce more structural variations for the person while retaining similar appearance. Models are trained in the COCO-Stuff dataset [22] and synthesized images are generated at a resolution of  $256 \times 256$  for both methods. See text for details.

objects (e.g., faces and cars) by the StyleGAN [8]<sup>1</sup>. Despite the big successes achieved by BigGANs and StyleGANs, learning generative models for layout-to-image entails more research. In addition to realness, generative models for layout-to-image need to tackle many spatial and semantic relationships among multiple objects (combinatorial in general). Specifically, learning layout-to-image requires addressing the problems of learning one-to-many mapping (*i.e.*, one layout covers many plausible realizations in image synthesis to preserve the intrinsic uncertainty), and of handling consistent multi-object generation (*e.g.*, occlusion handling for overlapped bounding boxes and uneven, especially long-tail distributions of objects). Because of those, it is difficult to capture underlying probability distributions defined in the solution space of layout-to-image.

In this paper, we further focus on *controllable image synthesis from reconfigurable layout and style*. As illustrated in Fig. 2, by controllable and reconfigurable, it means a generative model is capable of (i) **Style Control** – the model can preserve the intrinsic one-to-many mapping from a given layout to multiple plausible images with sufficiently different structural and appearance styles (*i.e.*, diversity), at

both image and object levels, and (ii) **Layout Control** – the model is also adaptive with respect to changes of layouts (*e.g.*, adding new objects), or perturbations of bounding boxes in a given layout, as well as the styles associated with the changes of spatial layouts. Prior arts on layout-to-image mainly focus on low resolution ( $64 \times 64$ ) [18], [20], except for the very recent Grid2Im method [19] which can synthesize images at a resolution of  $256 \times 256$ . Main drawbacks of existing methods are in two-fold: the diversity of styles in generated images is not sufficiently high to preserve the intrinsic one-to-many mapping for a given layout, and the controllability of styles along consecutive layout changes is often not sufficiently strong, as illustrated in Fig. 2. We aim to address these issues in this paper. Specifically, as we shall elaborate, the proposed method takes a step forward to address the problem of layout-to-image by learning layout-to-mask-to-image, while leveraging some of the best practice developed in state-of-the-art cGANs [6], BigGANs [5], StyleGANs [8] and Mask R-CNNs [25].

## 1.2 Method Overview

To learn controllable image synthesis from reconfigurable layouts and style codes, we build on Generative Adversarial Networks (GANs) [1] and present a *LayOut- and Style-based* architecture and learning paradigm for GANs. We

1. We view the StyleGAN as an implicitly conditional image synthesis framework since only one category is usually handled in training

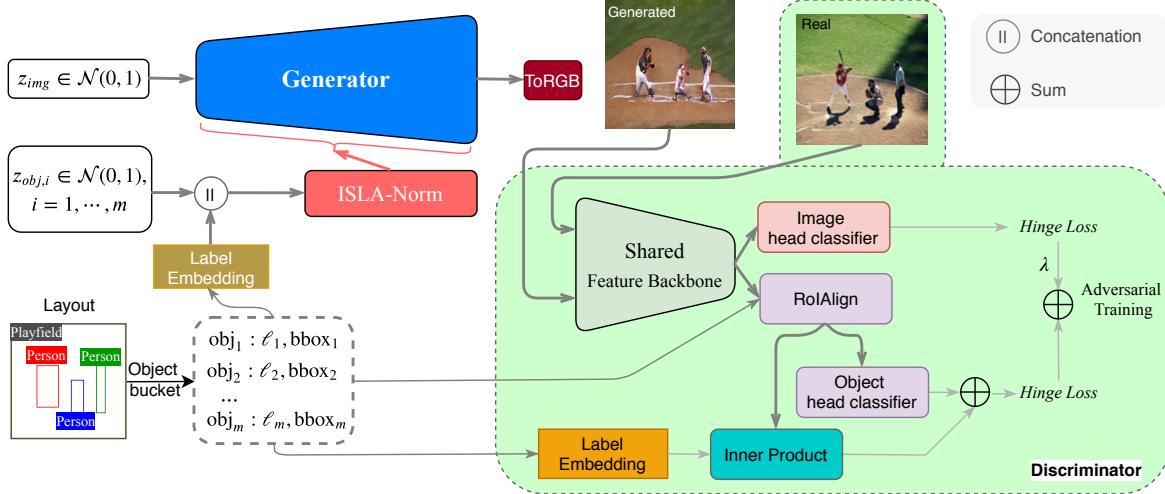


Fig. 3. Illustration of the network architecture of the proposed method. We build on GANs. Neural architectures of both the generator and discriminator use ResNets [24]. “ToRGB” is a simple module converting the final feature map in the generator to RGB images. The RoIAlign operation [25] is widely used in two-stage region-based ConvNets for object detection and semantic segmentation. Our proposed ISLA-Norm and detailed specifications of the generator are explained in Fig. 4. The discriminator is elaborated in Fig. 5. Note that the discriminator in the green box are not needed after training. See text for details. Best viewed in color.

termed the proposed method LostGAN in our previous conference paper presented at ICCV 2019, entitled *Image Synthesis from Reconfigurable Layout and Style* [26]. We shall call the conference version LostGAN-V1 and the updated model LostGAN-V2 in this paper. We first give an overview of our LostGAN and then summarize the changes of LostGAN-V2.

**The proposed LostGAN addresses the layout-to-image problem by learning layout-to-mask-to-image.** To account for the gap between bounding boxes in a layout and underlying object shapes, learning layout-to-mask is an intuitive and straightforward intermediate step to induce finer-grained style control of objects in a synthesized image, which also helps decouple learning of object geometry and learning of object appearance. Layout-to-mask itself is a relatively easier task than direct layout-to-image since object appearance are ignored. In the meanwhile, motivated by the impressive recent progress on conditional image synthesis from semantic label maps [11], [14], [15], it also makes sense to integrate layout-to-mask. If a reasonably good mask can be inferred for an input layout, learning mask-to-image can then leverage the best practice in conditional image synthesis from semantic label maps. A naïve approach is to develop two-stage generators, which may provide less effective solutions. Instead, we present a joint learning paradigm (*i.e.*, using a single generator). Fig. 3 illustrates the overall workflow of the proposed LostGAN. Fig. 4 illustrates the joint learning of layout-to-mask-to-image.

**The generator** has three inputs as commonly used in generative learning of layout-to-image: (i) a spatial layout,  $L$  consisting of a number of object bounding boxes in an image lattice, (ii) a latent vector,  $z_{img}$  for style control at the image level, and (iii) a bucket of latent vectors,  $z_{obj,i}$ 's, each of which is used for style control of an object instance. The latent vectors are randomly sampled from the standard multivariate Gaussian distribution. The generator takes the image latent vector as its direct input for overall style control, while utilizing a novel feature normalization scheme for object-level style control, which takes as input the concatenation of

the bucket of object latent vectors and the label embedding of objects in the layout. The object latent codes are involved in each stage of the generator for better style control, similar in spirit to the StyleGAN [8]. The objective of the generator is to capture the underlying probability distribution,  $p(I^{syn}|L, z_{img}, z_{obj_1}, \dots, z_{obj_m})$ . While straightforward for synthesizing images (using a single phase of forward computation), the generator involves a challenging inference step entailed in estimating the model parameters, that is to compute the latent codes for a real image  $I^{real}$  by sampling the posterior distribution,  $p(z_{img}, z_{obj_1}, \dots, z_{obj_m}|I^{real}, L)$ . To get around the difficulty of the posterior inference, GANs utilize an adversarial training paradigm by introducing an extra discriminator. Detailed specifications of the generator are shown in Fig. 4.

**The discriminator** has two inputs: a given image, either synthesized or real, and the corresponding spatial layout. It consists of three components: (i) a feature backbone extracting features from the input image, (ii) an image head classifier computing the image realness score based on the extracted features (the higher the score is, the more real an image is), and (iii) an object head classifier computing the realness score for each object instance. The features for an object instance is computed by the RoIAlign operation [25] using the given layout. Motivated by the projection-based cGANs [6] and the practice in the BigGAN [5], a label projection-based score is added to the realness score of each object instance. Detailed specifications of the generator are shown in Fig. 5.

**The loss function** consists of both image and object adversarial hinge loss terms [4], [6], [27], [28] (balanced by a trade-off parameter,  $\lambda$ ). The hinge loss aims to push the realness score of a synthesized image sufficiently away from that of a real image by a predefined margin. Under the two-player minmax game setting of GANs, the hinge loss works better to enforce both the generator and the discriminator more aggressive, leading to synthesized images of higher fidelity.

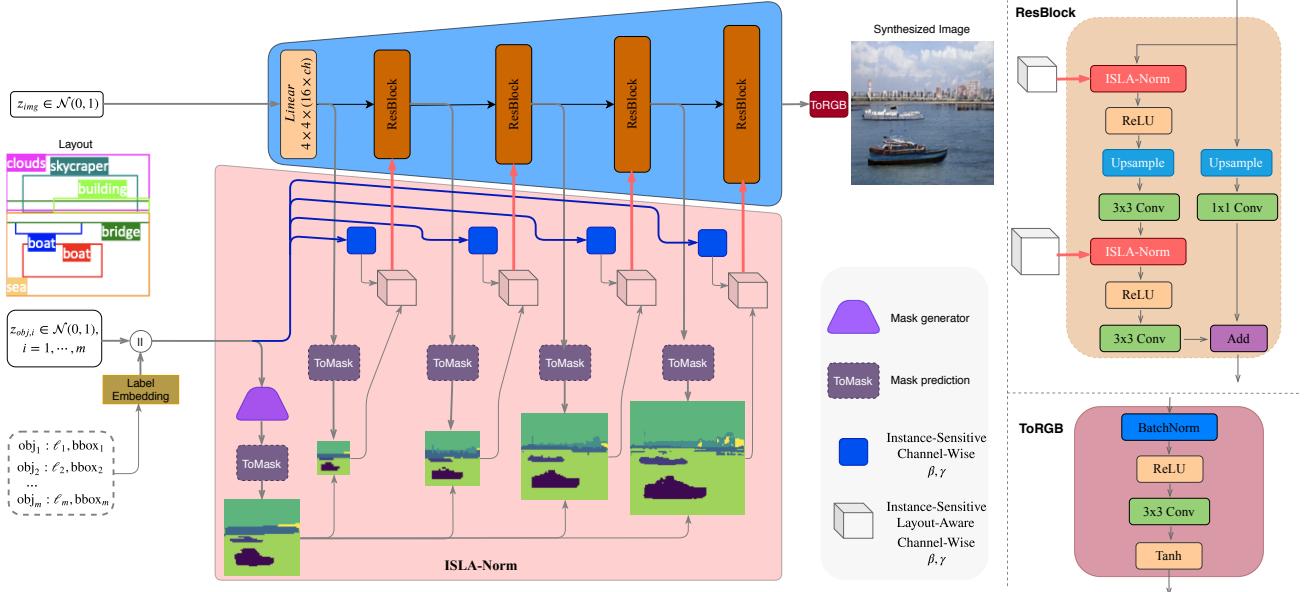


Fig. 4. Illustration of our proposed Instance-Sensitive and Layout-Aware Feature Normalization (ISLA-Norm) for the generator (*left*) and its deployment in a Residual building block (*right-top*). The *right-bottom* illustrates the “ToRGB” module. Our proposed ISLA-Norm realizes the learning of layout-to-mask-to-image for the generative learning problem of layout-to-image. The masks inferred on-the-fly enrich the outputs of our proposed model, going beyond traditional image synthesis and leading to joint image and label map synthesis. Note that the masks shown here are computed by a trained model for the given layout. During early stages of the training, the masks are much noisier. See text for details.

**The Instance-Sensitive and Layout-Aware Feature Normalization (ISLA-Norm)** scheme is presented to realize the proposed layout-to-mask-to-image pipeline in our LostGAN. Fig. 4 illustrates the proposed ISLA-Norm. As a feature normalization scheme, it consists of two components: feature standardization, and feature recalibration by learning an affine transformation. The former is done as the BatchNorm [29] in which channel-wise mean and standard deviation are computed in a mini-batch. The latter is different from BatchNorm. Unlike BatchNorm in which channel-wise affine transformation parameters,  $\beta$  (for re-shifting) and  $\gamma$  (for re-scaling) are learned as model parameters and shared across spatial dimensions by all instances, in our ISLA-Norm, we first learn object instance-sensitive channel-wise affine transformations from the concatenation of object label embedding and object style latent vectors, as shown by the arrows in blue in Fig. 4, similar in spirit to the Adaptive Instance Normalization (AdaIN) used in StyleGANs [8] and the projection-based conditional BatchNorm used in cGANs [5]. We also learn the mask for an input layout in two pathways: one pathway learns the mask from the concatenation of object label embedding and object style latent vectors, and the other learns the mask from feature maps at different stage in the generator. A learnable weighted sum of the two masks are used as the inferred mask at a stage in the generator. Then, to obtain fine-grained spatially-distributed multi-object style control for an input layout, we place the object instance-sensitive channel-wise affine transformations in the learned mask, leading to the instance-sensitive and layout-aware affine transformations for feature recalibration in the generator, as illustrated by the light-grey cube in Fig. 4.

**Summary of Changes.** Compared to LostGAN-V1, the main changes of LostGAN-V2 are as follows.

- The ISLA-Norm is extended by integrating masks learned from feature maps at different stages in the generator.
- The experiments are significantly extended by training models at higher resolutions and by comparing with the prior arts including the Grid2Im [19] and the GauGAN [15].
- The paper is thoroughly rewritten with much more details on different aspects of the LostGAN and on the experimental settings, together with new figures of the model.
- Several ablation studies are added to analyze the proposed LostGAN and ISLA-Norm.

Our source code and pretrained models have been made publicly available at <https://github.com/iVMCL/LostGANs>.

### 1.3 Related Work

Generative models have been studied widely in recent years such as Autoregressive models, Variational Autoencoders (VAEs) and Generative Adversarial Networks (GANs). For image generation, Autoregressive models such as PixelRNN [30] and PixelCNN [31] synthesize images pixel by pixel based on conditional distribution over pixels. VAEs [32], [33] jointly train an encoder and decoder where the former maps images into latent distribution and the latter generates images based on the latent distribution. GANs [1] are able to synthesize realistic and high resolution images under various settings, including both unconditional [2], [7], [8], [34] and conditional tasks [5], [6], [9]. Typically, a GAN consists of a Generator that produces realistic fake images from input (*e.g.*, random noise) and a Discriminator that distinguishes generated images from real ones. More recently, a unified divergence triangle framework is proposed for joint training of generator model, energy-based model, and inference model for generative tasks [35].

**Conditional Image Synthesis.** Conditional Image synthesis takes additional information (*i.e.*, class information [3], [5], [6], [9], [36], source image [10], [12], [13], [37], text description [16], [17], [38], [39], scene graph [18], [40], etc) as input. How to feed conditional information to a GAN model has been studied in various ways. In [9], [16], the conditional information are encoded to a vector, concatenated with noise, and then passed to the generator. In [6], projection-based methods, which incorporate conditional information to the discriminator by inner product between feature and learned class embedding, effectively improve the quality of class conditional image generation. In [16], [17], [41], conditional information is utilized by the discriminator using simple concatenation with the input or intermediate feature maps. In [5], [15], [42], [43], conditional information is provided to the generator by conditional gains and bias in the Batch Normalization [29] layers. The concurrent work, GauGAN [15] learns spatially adaptive normalization parameters from annotated semantic masks, while our proposed ISLA-Norm learns from coarse layout information.

**Image Synthesis from Layout.** Image synthesis from coarse layout, which usually used as intermediate representation for other conditional image synthesis tasks, has been studied in the very recent literature and proven a difficult task. In [18], [23], [39], [44], layout and object information are utilized in text to image or scene graph to image generation. In [44], [45], locations of multiple objects are controlled in text-to-image generation by adding an extra object pathway in both the generator and discriminator. In [18], [23], [39], a two-step approach is used in image synthesis: generating the semantic layout (class label, bounding boxes and semantic mask) from a text description or a scene graph, and synthesizing images conditioned on the predicted semantic layout and text description (if present). However, in [19], [23], [39], pixel-level instance segmentation annotations are needed in training, while our proposed method does not require pixel-level annotations and can learn semantic masks in a weakly-supervised manner. The layout-to-image task was first studied in [20] at the resolution of  $64 \times 64$ , which uses a variational autoencoders based network, together with long-short term memory (LSTM), for object feature fusion. In [40], an external memory bank is introduced, consisting of objects cropped from real images in training, which are retrieved and pasted in generating images from layouts at the resolution of  $64 \times 64$ .

## 1.4 Our Contributions

This paper makes the following main contributions to the field of conditional image synthesis.

- It presents a layout- and style-based architecture for GANs (termed LostGANs) which address the problem of layout-to-image by learning layout-to-mask-to-image and realize controllable image synthesis from reconfigurable layouts and styles. The outputs of our LostGANs cover both image and semantic mask synthesis.
- It presents an object instance-sensitive and layout-aware feature normalization scheme (termed ISLA-Norm) which explicitly and jointly accounts for the learning of layout-to-mask and the learning of spatially-distributed affine

transformations for feature recalibration at an object mask level.

- It can synthesize images at a resolution of up to  $512 \times 512$  and shows state-of-the-art performance in terms of the Inception Score [46], the Fréchet Inception Distance [47], the Diversity Score base on the LPIPS metric [48] and the classification accuracy [49] on two widely used datasets, the COCO-Stuff [22] and the Visual Genome [50].

## 1.5 Paper Organization

In the remainder of this paper, Section 2 presents the problem formulation of layout-to-image and technical details of our proposed LostGAN and ISLA-Norm. Section 3 shows the experimental settings, quantitative and qualitative results, together with an ablation study. Section 4 concludes this paper.

## 2 THE PROPOSED METHOD

In this section, we first define the problem and then present details of our LostGAN and ISLA-Norm.

### 2.1 Problem Formulation

Denote by  $\Lambda$  an image lattice (*e.g.*,  $256 \times 256$ ) and by  $I$  an image defined on the lattice. Let  $L = \{(\ell_i, bbox_i)\}_{i=1}^m$  be a layout consisting of  $n$  labeled bounding boxes, where label  $\ell_i \in \mathcal{C}$  (*e.g.*,  $|\mathcal{C}| = 171$  in the COCO-Stuff dataset [22]), and a bounding box  $bbox_i \subseteq \Lambda$ . Different bounding boxes may overlap and thus have undetermined partial-order of occlusions. Let  $z_{img}$  be the latent code controlling image style and  $z_{obj_i}$  the latent code controlling object instance style for  $(\ell_i, bbox_i)$  (*e.g.*, the latent codes are randomly sampled from the standard Gaussian distribution,  $\mathcal{N}(0, 1)$  under the i.i.d. setting). Denote by  $Z_{obj} = \{z_{obj_i}\}_{i=1}^m$  the set of object instance style latent codes.

Image synthesis from layout and style is to learn a generation function capable of synthesizing an image for a given input  $(L, z_{img}, Z_{obj})$ ,

$$I^{syn} = \mathcal{G}(L, z_{img}, Z_{obj}; \Theta_{\mathcal{G}}) \quad (1)$$

where  $\Theta_{\mathcal{G}}$  represents the parameters of the generation function. In general, a generator model  $\mathcal{G}(\cdot)$  is expected to capture the underlying conditional data distribution  $p(I|L, z_{img}, Z_{obj}; \Theta_{\mathcal{G}})$  in the high-dimensional space.

**Reconfigurability of a generator model  $\mathcal{G}(\cdot)$ .** We are interested in three aspects in this paper:

- *Image style reconfiguration:* If we fix the layout  $L$ , is the generator  $\mathcal{G}(\cdot)$  capable of synthesizing images with different styles for different  $(z_{img}, Z_{obj})$ , while retaining the object configuration conditioned on  $L$ ?
- *Object style reconfiguration:* If we fix the layout  $L$ , the image style  $z_{img}$  and all the object styles  $Z_{obj}$  except for  $z_{obj_i}$ , is the generator  $\mathcal{G}(\cdot)$  capable of generating consistent images with different styles for the object  $(\ell_i, bbox_i)$  using different  $z_{obj_i}$ , while retaining the object configuration conditioned on  $L$  and the styles of the remaining objects?
- *Layout reconfiguration:* Given an input tuple  $(L, z_{img}, Z_{obj})$ , is the generator  $\mathcal{G}(\cdot)$  capable of generating consistent images for different  $(L^+, z_{img}, Z_{obj}^+)$ 's, where we can add a new object to  $L^+$  or just change the location

and/or label of an existing bounding box? When a new object is added, we also sample a new  $z_{obj}$  to add in  $Z_{obj}^+$ . When only the bounding box location changes, we keep all latent codes unchanged (*i.e.*,  $Z_{obj}^+ = Z_{obj}$ ).

It is a challenging problem to address the three aspects by learning a single generator model  $\mathcal{G}(\cdot)$ . Intuitively, it might be even difficult for well-trained artistic people to do so at scale (*e.g.*, handling the 171 categories in the COCO-Stuff dataset). Due to the complexity that the generator model (Eqn. 1) needs to handle, it is often implemented using expressive and over-parameterized deep neural networks (DNNs). It is also well-known that training a DNN-based generator model individually is a extremely difficult task (due to the difficulty of sampling the posterior). Adversarial training is a popular workaround and GANs [1] are widely used in practice, which are formulated under a two-player minmax game setting.

## 2.2 The LostGAN

In this section, we present the technical details of our LostGAN.

### 2.2.1 The Generator

As illustrated in Fig. 4, the generator  $\mathcal{G}(\cdot)$  consists of a linear full-connected (FC) layer, followed by a number of residual building blocks (ResBlocks) [24] depending on the target resolution of image synthesis, and a “ToRGB” module outputting a synthesized image.

- The image style latent code  $z_{img} \in \mathbb{R}^{d_{img}}$  is a  $d_{img}$ -dim vector (*e.g.*,  $d_{img} = 128$  in our experiments). The linear FC layer projects  $z_{img}$  to a  $4 \times 4 \times (16 \times ch)$  dimensional vector. The projected vector is then reshaped as a tensor of dimensions  $(4, 4, 16 \times ch)$  (representing height, width and channels), where  $ch$  is a hyperparameter to control the model complexity (*e.g.*,  $ch = 64$  used in our experiments).
- The right-top of Fig. 4 shows the detail of a ResBlock. The ResBlock uses the basic block design in ResNets [24], as adopted in the projection-based cGAN [6] and the BigGAN [5]. Each ResBlock upsamples its input once by a factor of 2. So, we will need  $B$  ResBlocks to generate images at a resolution of  $4^{B-1} \times 4^{B-1}$  (*e.g.*,  $B = 4$  for the resolution of  $64 \times 64$ ). Each ResBlock also reduces the number of feature channels of its input by a factor of 2.
- The final “ToRGB” module consists of a BatchNorm [29] layer, a ReLU layer, a  $3 \times 3$  convolution layer with output channels being 3, and a Tanh layer, as illustrated in the right-bottom of Fig. 4.

### 2.2.2 The ISLA-Norm

The are two ISLA-Norm modules in a ResBlock. Denote by  $x$  an input 4D feature map of ISLA-Norm, and  $x_{n,c,h,w}$  the feature response at a position  $(n, c, h, w)$  (using the convention order of axes for batch, channel, and spatial height and width). We have  $n \in [0, N-1]$ ,  $c \in [0, C-1]$ ,  $h \in [0, H-1]$ ,  $w \in [0, W-1]$ , where  $N$  is the mini-batch size or the accumulated size of synchronized mini-batches, and  $C, H, W$  depend on the stage of a ResBlock.

**Feature Standardization.** Our ISLA-Norm first computes the channel-wise mean and standard deviation as done in

the BatchNorm [29]. In training, ISLA-Norm first normalizes  $x_{n,c,h,w}$  by,

$$\hat{x}_{n,c,h,w} = \frac{x_{n,c,h,w} - \mu_c}{\sigma_c}, \quad (2)$$

where the channel-wise batch mean  $\mu_c = \frac{1}{N \cdot H \cdot W} \sum_{n,h,w} x_{n,c,h,w}$  and standard deviation  $\sigma_c = \sqrt{\frac{1}{N \cdot H \cdot W} \sum_{n,h,w} (x_{n,c,h,w} - \mu_c)^2 + \epsilon}$  ( $\epsilon$  is a small positive constant for numeric stability).

**Feature Recalibration.** In the vanilla BatchNorm [29], the recalibration is done by learning channel-wise affine transformations, consisting of the re-scaling paramter,  $\gamma_c$ 's and the re-shifting parameters,  $\beta_c$ 's. We have,

$$\tilde{x}_{n,c,h,w}^{BN} = \gamma_c \cdot \hat{x}_{n,c,h,w} + \beta_c. \quad (3)$$

In our ISLA-Norm, we will learn instance-sensitive and layout-aware affine transformation parameters,  $\gamma_{n,c,h,w}$ 's and  $\beta_{n,c,h,w}$ 's, and we have,

$$\tilde{x}_{n,c,h,w} = \gamma_{n,c,h,w} \cdot \hat{x}_{n,c,h,w} + \beta_{n,c,h,w}. \quad (4)$$

**Computing  $\gamma_{n,c,h,w}$  and  $\beta_{n,c,h,w}$ .** Without loss of generality, we show how to compute the gamma and beta parameters for one sample, *i.e.*,  $\gamma_{c,h,w}$  and  $\beta_{c,h,w}$ . As shown in the left of Fig. 4, we have the following six components.

*i) Label Embedding.* We use one-hot label vector for the  $m$  object instances in a layout  $L$ , which results in a one-hot label matrix, denoted by  $Y$ , of the size  $m \times d_\ell$ , where  $d_\ell$  is the number of object categories (*e.g.*,  $d_\ell = 171$  in the COCO-Stuff dataset). Label embedding is to learn a  $d_\ell \times d_e$  embedding matrix, denoted by  $W$ , to compute the vectorized representation for labels,

$$Y = Y \cdot W, \quad (5)$$

where  $Y$  is a  $m \times d_e$  matrix and  $d_e$  represents the embedding dimension (*e.g.*,  $d_e = 128$  in our experiments).

*ii) Joint Label and Style Encoding.* We sample from the standard Gaussian distribution the object style latent codes  $Z_{obj}$  which is a  $m \times d_{obj}$  noise matrix (*e.g.*,  $d_{obj} = 128$  the same as  $z_{img}$  in our experiments). We concatenate the label-to-vector matrix  $Y$  and the object latent style matrix as the joint label and style encoding,

$$S = (Y, Z_{obj}), \quad (6)$$

which is a  $m \times (d_e + d_{obj})$  matrix. So, the object instance style depends on both the label embedding (semantics) and i.i.d. latent codes (accounting for style variations).

*iii) Mask Generation from the joint label and style encoding S.* We first generate an mask for each object instance in a layout  $L$  at some predefined size,  $s \times s$  (*e.g.*,  $s = 32$ ) individually. Then, we resize the generated masks to the sizes of corresponding bounding boxes at a ResBlock stage in the generator.

• The mask generation process consists of two components: one is a simplified generator model (the small trapezoid in purple in Fig. 4), and the other is a simple “ToMask” operation. The former composes a linear FC layer projecting  $S$  to a tensor with a resolution of  $4 \times 4$  (after reshaping), and a few stages of Conv3x3+BatchNorm+ReLU (where Conv3x3 refers to a convolution operation with a kernel of spatial dimensions,  $3 \times 3$ ). The “ToMask” operation is

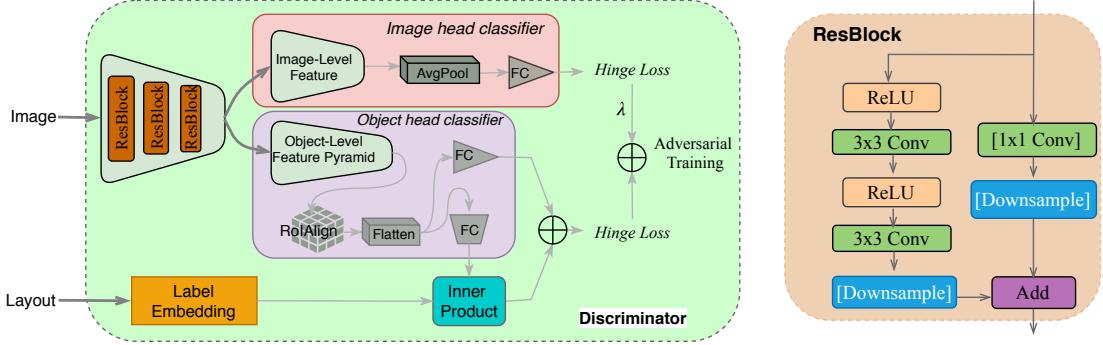


Fig. 5. Illustration of the discriminator (*left*). The shared feature backbone, the image-level feature backbone and the object-level feature pyramid use ResBlocks (*right*), and each consists of a number of ResBlocks depending on the target resolution of layout-to-image (e.g.,  $128 \times 128$  or  $256 \times 256$ ). In the ResBlock, “[op]” means an operation is optional subject to the settings. The object-level feature pyramid is used for placing object instances of different sizes at different feature layers (e.g., smaller bounding boxes placed at lower feature layers as done in the FPN [51]), such that the RoIAlign operation is meaningful. “FC” represents a fully-connected layer (with either a scalar output shown by a grey triangle or a vector output shown by a grey trapezoid). “AvgPool” represents a global channel-wise average pooling over the spatial dimensions. See text for details.

then implemented by Conv3x3+Sigmoid, whose output is  $m \times s \times s$ , representing a  $s \times s$  mask for each of the  $m$  object instances. Based on the mask size  $s$ , we also upsample the feature map after a Conv3x3+BatchNorm+ReLU by a factor of 2 for a number of stages as needed.

- After resizing and reassembling the generated mask for a ResBlock stage, we obtain a mask tensor of dimensions  $(m, H, W)$ , denoted by  $\mathbb{M}_S$ , each slice of which has zeros outside the corresponding bounding box,  $bbox_i$ . For the visualization purpose (e.g., those shown in Fig. 4), when reassembling the resized object masks, we use arg max across the  $m$  channels of  $\mathbb{M}_S$  to assign the label index for a pixel occupied by more than one objects due to occlusions.

*iv) Mask Generation from the feature maps in a generator.* For a ResBlock stage, we learn a mask from its input feature map using a simple “ToMask” operation implemented by Conv3x3+Sigmoid, where the out channel of the Conv3  $\times$  3 kernel is  $d_\ell$  (*i.e.*, the number of categories in a dataset). The mask is represented by a tensor of sizes  $(d_\ell, H, W)$ . We clip the mask based on the layout by only keeping values unchanged within the bounding boxes of the object instances in a layout (and zeroing out the remainder). Denote by  $\mathbb{M}_F(L)$  the mask tensor of sizes  $(m, H, W)$  after the clipping (omitting the index for a ResBlock in the generator without loss of generality). For the second ISLA-Norm in a ResBlock (the right-top of Fig. 4), we just upsample  $\mathbb{M}_F(L)$  by a factor of 2 for simplicity.

*v) Object instance-sensitive channel-wise affine transformations learned from the joint label and style encoding  $\mathbb{S}$ .* We adopt a linear projection with a learnable  $(d_e + d_{obj}) \times 2C$  projection matrix  $\mathcal{A}$ , where  $C$  is the number of channels, and we have,

$$\mathcal{T} = \mathbb{S} \cdot \mathcal{A}, \quad (7)$$

which is a matrix of sizes  $(m, 2C)$ . Let  $\mathcal{T}_\beta$  and  $\mathcal{T}_\gamma$  be the column-wise first and second half of  $\mathcal{T}$ . We unsqueeze both  $\mathcal{T}_\beta$  and  $\mathcal{T}_\gamma$  to the size of  $(m, C, H, W)$  by replicating values across the spatial dimensions. Learning the affine transformations in this way leads to stronger style control of our LostGAN than other layout-to-image methods, since the style latent codes get involved in every stage of the

generator, rather than being used as input only to the first stage of the generator in other layout-to-image methods.

*vi) Computing ISLA  $\gamma_{c,h,w}$  and  $\beta_{c,h,w}$ .* We first unsqueeze the two masks,  $\mathbb{M}_S$  and  $\mathbb{M}_F(L)$ , to the sizes  $(m, C, H, W)$  by replicating  $C$  channels. Then, we have,

$$\gamma_{c,h,w} = \frac{1}{M_{h,w}} \sum_{i=1}^m \mathbb{M}(i, c, h, w) \times \mathcal{T}_\gamma(i, c, h, w), \quad (8)$$

$$\beta_{c,h,w} = \frac{1}{M_{h,w}} \sum_{i=1}^m \mathbb{M}(i, c, h, w) \times \mathcal{T}_\beta(i, c, h, w), \quad (9)$$

where  $\mathbb{M}(\cdot) = [(1 - \alpha) \cdot \mathbb{M}_S + \alpha \cdot \mathbb{M}_F(L)](\cdot)$  with  $\alpha$  being a learnable weight to balance the two masks, and  $M_{h,w} = \sum_{i=1}^m \mathbb{M}(i, 0, h, w)$  if the pixel  $(h, w)$  is occupied by multiple object bounding boxes ( $M_{h,w}$  is then to normalize the sum of mask values), otherwise  $M_{h,w} = 1$ . We note that unlike the arg max operation that we used for the overlapping regions in visualizing the learned masks, the average is used in computing the spatially-distributed affine transformations.

**Handling Background.** To account for the situation in which all object instances do not occupy the entire image lattice (e.g., in the VG dataset [50]), we introduce a background class  $\ell_0$  with  $bbox_0 = \Lambda$ .

### 2.2.3 The Discriminator

As shown in Fig. 5, our discriminator consists of three components: a shared ResNet-based feature backbone, an image head classifier and an object head classifier.

- *The shared ResNet-based feature backbone* consists of a number of basic ResBlocks [24] with the BatchNorm modules removed (as commonly adopted in prior arts such as the SNGAN [4] and the BigGAN [5]). The number of ResBlocks is subject to the target resolution of layout-to-image.
- *The image head classifier* consists of an image-level feature backbone, a global average pooling layer, and an one-output FC layer. The image-level feature backbone includes a number of ResBlocks. The output of the FC layer is a scalar realness score, denoted by  $p_{img}$ .
- *The object head classifier* is a simplified version of R-CNN based object detector [25], [52] with bounding boxes being

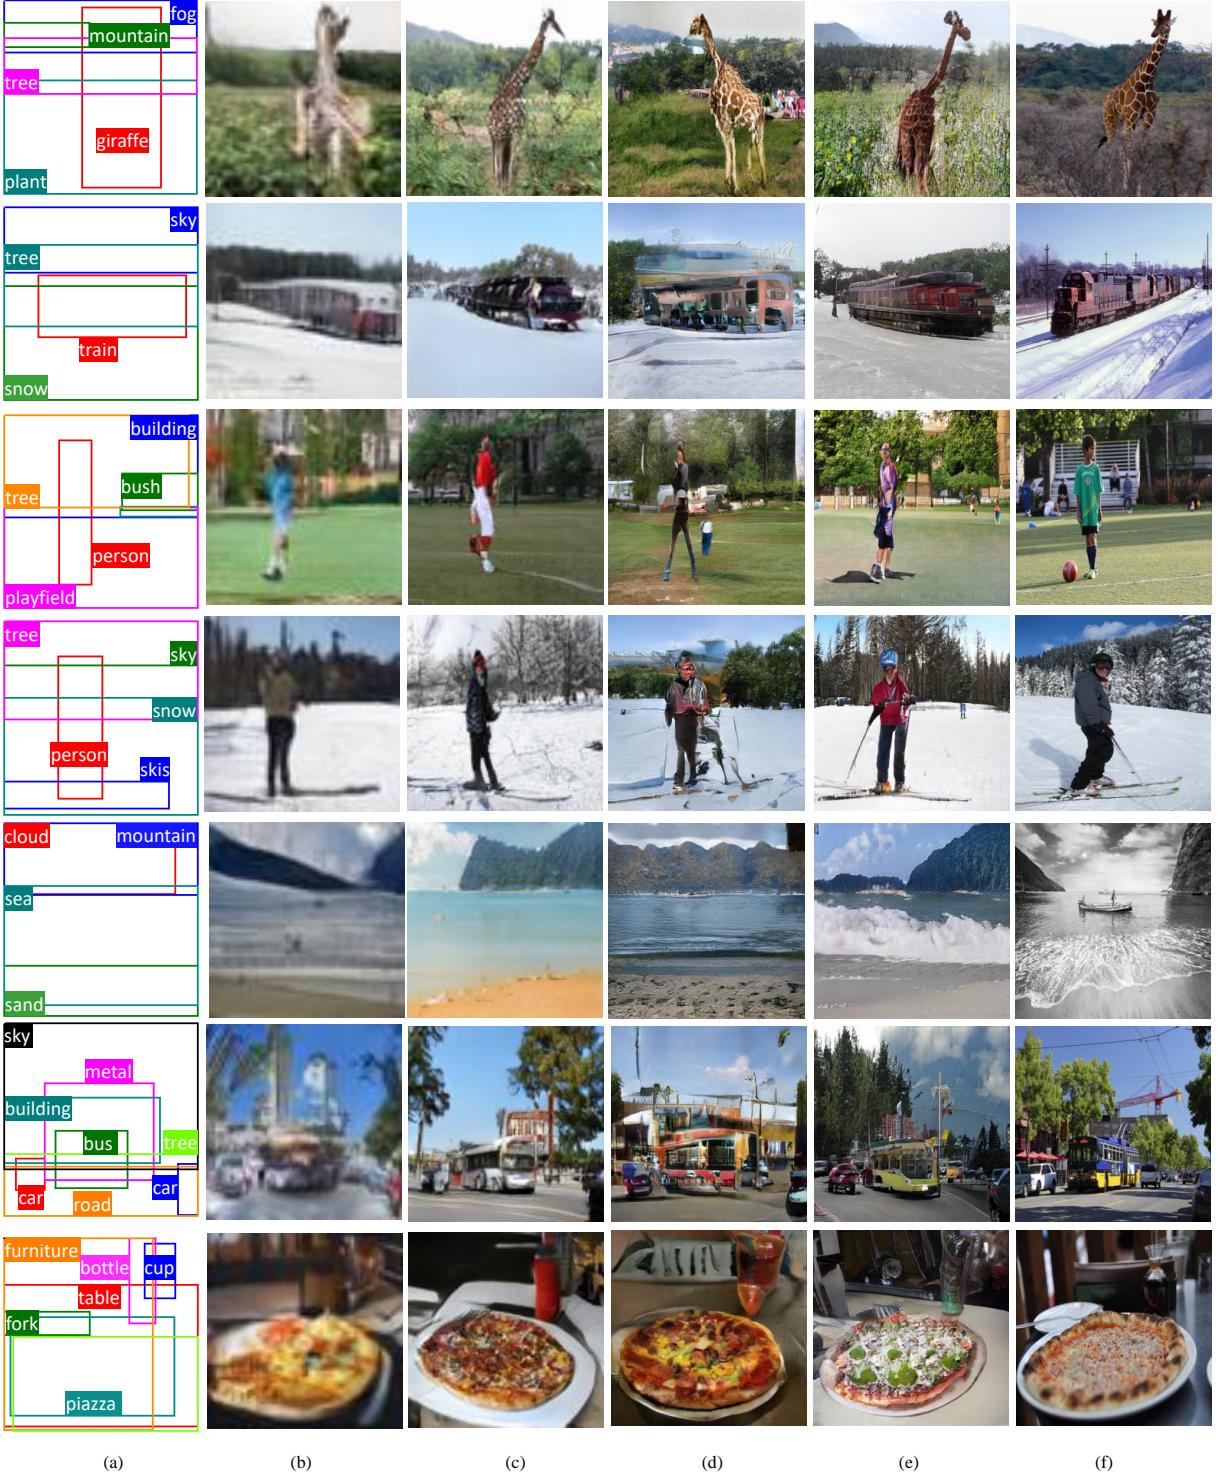


Fig. 6. Generated samples from given layouts in COCO-Stuff by different models. (a) Input Layout, and Images generated by (b) Layout2Im [20] 64 x 64, (c) our LostGAN-V1 128 x 128, (d) Grid2Im [19] 256 x 256, (e) our LostGAN-V2 256 x 256, and (f) Ground Truth.

given in the input layout. It consists of an object-level feature pyramid, a RoIAlign layer [25], and two FC layers. The one-output FC layer computes a scalar realness score for each object instance, denoted by  $p_{obj_i}^r$  ( $i \in [1, m]$ ). To encode object label semantic information, we also compute a label projection score [5], [6] for each object instance, denoted by  $p_{obj_i}^\ell$ , which is the inner product between the label embedding (calculated in the same as

Eqn. 5) and the linear projection (using a FC layer) of the RoIAlign feature vector. The overall score of an object instance is the sum:  $p_{obj_i} = p_{obj_i}^r + p_{obj_i}^\ell$ .

In sum, denote by  $\mathcal{D}(\cdot; \Theta_D)$  the discriminator with parameters  $\Theta_D$ . Given an image  $I$  (real or synthesized) and a layout  $L$ , the discriminator computes a list of scores,

$$(p_{img}, p_{obj_1}, \dots, p_{obj_m}) = \mathcal{D}(I, L; \Theta_D) \quad (10)$$

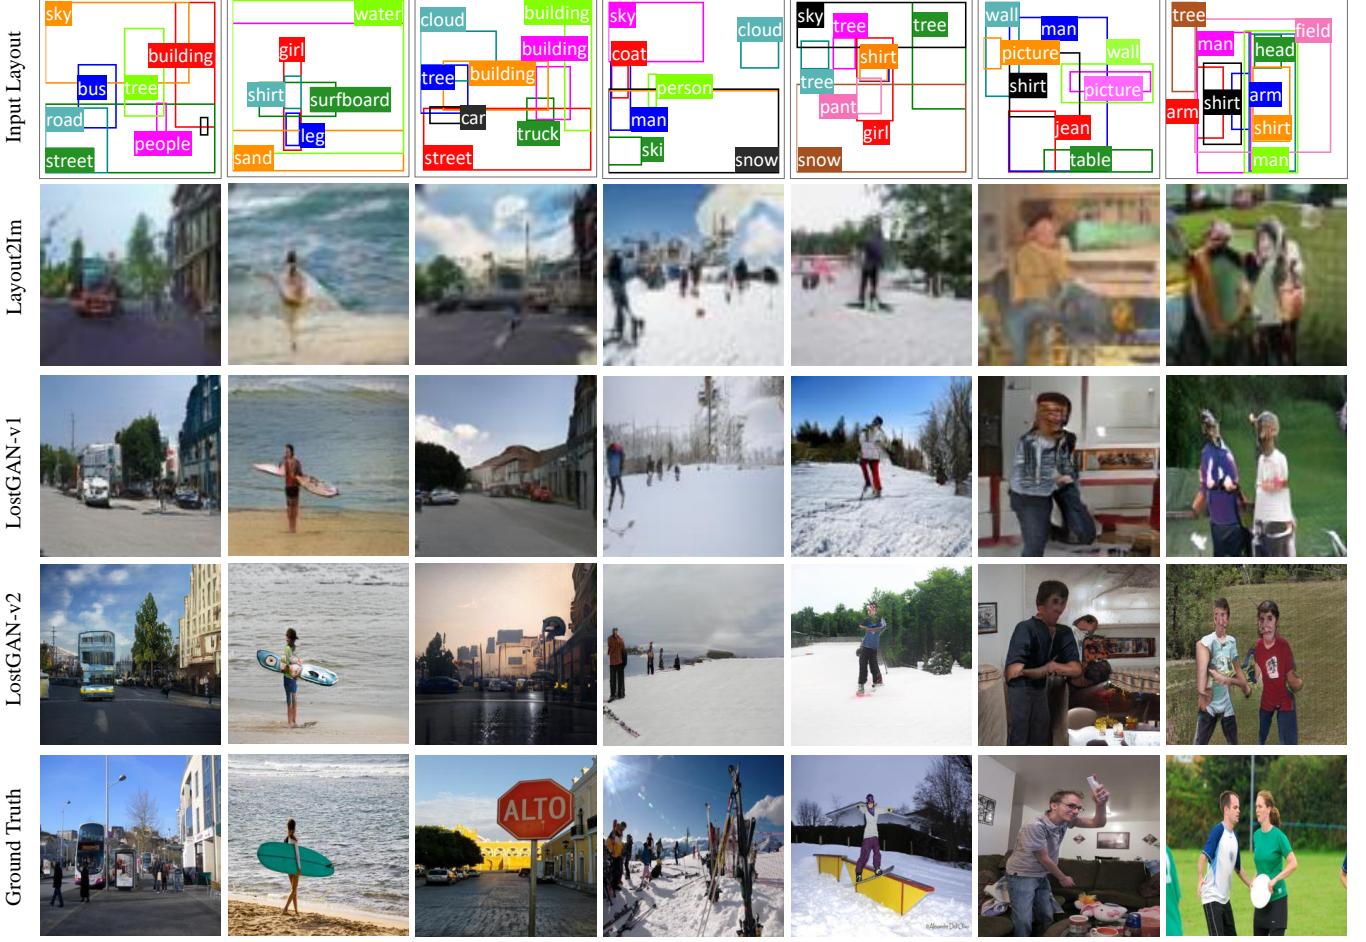


Fig. 7. Generated samples from given layouts on VG by different models. From top to bottom shows Input Layout, and Images generated by Layout2Im [20]  $64 \times 64$ , our LostGAN-V1  $128 \times 128$ , our LostGAN-V2  $256 \times 256$ , and Ground Truth.

#### 2.2.4 The Loss Functions

Under the mini-batch based SGD framework, for the generator, the loss function of  $\Theta_G$  is defined by,

$$\mathcal{L}(\Theta_G | \Theta_D) = - \sum_{(L, I^{syn}, I^{gt}) \in \mathbb{B}} [P_{\mathcal{D}(I^{syn}, L; \Theta_D)} - ||I^{syn} - I^{gt}||_1 - ||F(I^{syn}) - F(I^{gt})||_1], \quad (11)$$

where  $\mathbb{B}$  represents a mini-batch,  $I^{syn}$  and  $I^{gt}$  represent a synthesized image (Eqn. 1) and the ground-truth image for the spatial layout  $L$ ,  $P_{\mathcal{D}(I, L; \Theta_D)} = \lambda \cdot p_{img} + \frac{1}{m} \sum_{i=1}^m p_{obj_i}$  with a trade-off parameter  $\lambda$  (0.1 used in our experiments), the second term in the right-hand side is the reconstruction loss, and the last term is the perceptual loss [53] which measure L1 difference between features,  $F(\cdot)$  of generated image and ground truth images by an ImageNet pretrained network such as the VGG network [54]. Minimizing  $\mathcal{L}(\Theta_G | \Theta_D)$  is trying to fool the discriminator by generating high fidelity images.

For the discriminator, we utilize the hinge version [27], [28] of the standard adversarial loss [1],

$$l_t(I, L) = \begin{cases} \max(0, 1 - p_t); & \text{if } I \text{ is a real image} \\ \max(0, 1 + p_t); & \text{if } I \text{ is a fake image} \end{cases} \quad (12)$$

where  $t \in \{img, obj_1, \dots, obj_m\}$ . In the hinge loss, no penalty will occur if the score of a real image (or a real

object instance) is greater than or equal to 1, and the score of a fake image (or a fake object instance) is less than or equal to -1. The hinge loss is more aggressive than the real vs fake binary classification in the vanilla GAN. The overall loss is,

$$l(I, L) = \lambda \cdot l_{img}(I, L) + \frac{1}{m} \sum_{i=1}^m l_{obj_i}(I, L). \quad (13)$$

The loss function of  $\Theta_D$  is defined by,

$$\mathcal{L}(\Theta_D | \Theta_G) = \sum_{(L, I^{syn}, I^{gt}) \in \mathbb{B}} [l(I^{gt}, L) + l(I^{syn}, L)], \quad (14)$$

where  $p(I, L)$  represents both the real and fake (synthesized by the generator) data. Minimizing  $\mathcal{L}(\Theta_D | \Theta_G)$  is trying to tell apart the real and fake images.

#### 2.2.5 Implementation Details

In implementation and training, we follow the practice used in [3], [5], [6]. Synchronized Batch Normalization [29], where batch statistics for feature standarization are computed over all devices, is adopted in our ISLA-Norm. The Spectral Normalization [4] of model parameters is also applied in both the Generator and the Discriminator to stabilize training. Parameters of the Generator and the Discriminator are initialized using the Orthogonal Initialization method [55]. The Adam optimizer [56] is used with  $\beta_1 = 0$  and  $\beta_2 = 0.999$ .

TABLE 1

Quantitative comparisons using the Inception Score (IS, higher is better), FID (lower is better) and Diversity Score (DS, higher is better) evaluation metrics in the COCO-Stuff [22] and VG [50] datasets. See text for details.

| Methods                          | IS↑                 |              | FID↓         |              | DS↑                |             |
|----------------------------------|---------------------|--------------|--------------|--------------|--------------------|-------------|
|                                  | COCO                | VG           | COCO         | VG           | COCO               | VG          |
| Real Images (64×64)              | 16.30 ± 0.40        | 13.90 ± 0.50 | -            | -            | -                  | -           |
| Real Images (128×128)            | 22.30 ± 0.50        | 20.50 ± 1.50 | -            | -            | -                  | -           |
| Real Images (256×256)            | 28.10 ± 1.60        | 28.60 ± 1.20 | -            | -            | -                  | -           |
| Real Images (512×512)            | 34.50 ± 1.70        | 34.20 ± 1.10 | -            | -            | -                  | -           |
| pix2pix [11] 64×64               | 3.50 ± 0.10         | 2.70 ± 0.02  | 121.97       | 142.86       | 0                  | 0           |
| sg2im (GT Layout) [18] 64×64     | 7.30 ± 0.10         | 6.30 ± 0.20  | 67.96        | 74.61        | 0.02 ± 0.01        | 0.15 ± 0.12 |
| Layout2Im [20] 64×64             | 9.10 ± 0.10         | 8.10 ± 0.10  | 44.19        | 39.68        | 0.15 ± 0.06        | 0.17 ± 0.09 |
| Layout2Im + OWA [20] 64×64       | 9.70 ± 0.10         | 8.00 ± 0.20  | 40.19        | 33.54        | 0.09 ± 0.05        | 0.09 ± 0.11 |
| LostGAN-V1 [26] 64×64            | 9.80 ± 0.20         | 8.70 ± 0.40  | 34.31        | 34.75        | 0.35 ± 0.09        | 0.34 ± 0.10 |
| Grid2Im [19] (GT Layout) 128×128 | 11.22 ± 0.15        | -            | 63.44        | -            | 0.28 ± 0.11        | -           |
| LostGAN-V1 [26] 128×128          | 13.80 ± 0.40        | 11.10 ± 0.60 | 29.65        | 29.36        | 0.40 ± 0.09        | 0.43 ± 0.09 |
| <b>LostGAN-V2</b> 128×128        | <b>14.21 ± 0.40</b> | 10.71 ± 0.26 | <b>24.76</b> | <b>29.00</b> | <b>0.45 ± 0.09</b> | 0.42 ± 0.09 |
| Grid2Im [19] (GT Layout) 256×256 | 15.23 ± 0.11        | -            | 65.95        | -            | 0.34 ± 0.13        | -           |
| <b>LostGAN-V2</b> 256×256        | <b>18.01 ± 0.50</b> | 14.10 ± 0.38 | <b>42.55</b> | 47.62        | <b>0.55 ± 0.09</b> | 0.53 ± 0.09 |
| <b>LostGAN-V2</b> 512×512        | 17.55 ± 0.23        | 14.42 ± 0.46 | 51.99        | 52.73        | 0.65 ± 0.11        | 0.61 ± 0.09 |

TABLE 2  
Comparisons of the CAS. See text for details.

| Methods                 | CAS↑  |       |
|-------------------------|-------|-------|
|                         | COCO  | VG    |
| Layout2Im [20] 64x64    | 27.32 | 23.25 |
| LostGAN-V1 [26] 64x64   | 28.81 | 27.50 |
| Grid2Im [19] 128x128    | 25.89 | -     |
| LostGAN-V1 [26] 128x128 | 30.68 | 28.85 |
| LostGAN-V2 128x128      | 31.98 | 29.35 |
| Grid2Im [19] 256x256    | 20.54 | -     |
| LostGAN-V2 256x256      | 30.33 | 28.81 |
| Real Images             | 51.04 | 48.07 |

The learning rate is set constant  $10^{-4}$  for both the Generator and the Discriminator. We use batch size of 128 based on our computing hardware resource. Training our LostGAN takes about 2-3 days, *e.g.*, either for the models generating  $128 \times 128$  images on 2 NVIDIA V100 GPUs, or for the models synthesizing  $256 \times 256$  images on 4 NVIDIA V100 GPUs.

### 3 EXPERIMENTS

We test our LostGAN in the COCO-Stuff dataset [22] and the Visual Genome (VG) dataset [50]. We evaluate LostGAN-V1 at two resolutions ( $64 \times 64$  and  $128 \times 128$ ) and LostGAN-V2 at three resolutions ( $128 \times 128$ ,  $256 \times 256$  and  $512 \times 512$ ). Our LostGAN-V2 obtains state-of-the-art performance.

#### 3.1 Datasets

The **COCO-Stuff** 2017 [22] augments the COCO dataset with pixel-level stuff annotations. The annotation contains 80 *thing* classes (person, car, *etc.*) and 91 *stuff* classes (sky, road, *etc.*) Following settings of [18], objects covering less than 2% of the image area are ignored, and we use images with 3 to 8 objects. For the **Visual Genome** (VG) dataset [50], we follow the settings of [18] to remove small and infrequent objects, which results in 62,565 images for training, 5,506 images for validation and 5,088 images for testing, with 3 to 30 objects from 178 categories in each image.

#### 3.2 Methods in Comparison

We compare with four prior arts: *i) The pix2pix* method [11] learns to map images between two domains. We reuse

the pix2pix results reported in the Layout2Im [20] in our comparisons, where a pix2pix model is trained to synthesize images from a feature map learned to encode the layout. The number of channels of the feature map is the number of categories (*e.g.*, 171 in COCO-Stuff). *ii) The scene graph to image (sg2im)* method [18] synthesizes images from input scene graphs with an intermediate scene-graph-to-layout module. We compare with sg2im using the ground-truth (GT) layouts. *iii) The Layout2Im* method [20] is the first to synthesize images directly from input layouts. These three methods have only been evaluated at the resolution of  $64 \times 64$ . *iv) The Grid2Im* method [19] extends the sg2im method, which has been tested at two resolutions,  $128 \times 128$  and  $256 \times 256$ , in the COCO-Stuff dataset only since ground-truth masks are needed in training. We also compare with Grid2Im using the GT layouts. Fig. 6 and 7 show examples of synthesized images by different methods.

#### 3.3 Evaluation Metrics

It remains a challenging problem to automatically evaluated DNN-based generator models in general. For layout-to-image, we adopt four state-of-the-art metrics as follows.

The *Inception Score* (IS) [46] uses an Inception V3 network pretrained on the ImageNet-1000 classification benchmark and computes a score (statistics) of the network’s outputs with  $N$  synthesized images  $I_i$ ’s of a generator model  $\mathcal{G}$ ,

$$IS(\mathcal{G}) \approx \exp\left\{\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^{1000} p(y=j|I_i) \log \frac{p(y=j|I_i)}{\hat{p}(y=j)}\right\}, \quad (15)$$

where  $\hat{p}(y=j) = \frac{1}{N} \sum_{i=1}^N p(y=j|I_i)$ . The IS aims to capture two desirable qualities of image synthesis: Synthesized image should contain clear and meaningful objects (subject to the ImageNet-1000 training datasets), and diverse images from all the different categories in ImageNet should be observed in synthesized images. So, *the larger the IS is, the better a generator model is*. Multiple runs are usually used to calculate the mean±std evaluation (*e.g.*, 5 runs are typically used). The IS does not leverage the statistics of real images.

The *Fréchet Inception Distance* (FID) [47] has been proposed to improve IS by incorporating statistics from real images. It also uses an ImageNet-pretrained Inception V3

network and computes the Frèchet distance [57] between two Gaussian distributions fitted to synthesized images and real images respectively. Denote by  $(\mu^0, \Sigma^0)$  and  $(\mu^1, \Sigma^1)$  the mean vector and the covariance matrix of the Gaussian distribution fitted on synthesized images and real images respectively. The Frèchet distance is defined by,

$$\text{FID}^2((\mu^0, \Sigma^0), (\mu^1, \Sigma^1)) = \|\mu^0 - \mu^1\|_2^2 + \text{Tr}(\Sigma^0 + \Sigma^1 - 2(\Sigma^0 \Sigma^1)^{1/2}). \quad (16)$$

So, the lower the FID is, the better a generator model is. Both the IS and FID do not explicitly measure the quality of one-to-many mapping in layout-to-image.

The *Diversity Score* (DS) aims to compare the perceptual similarity in a DNN feature space between two images,  $I_1$  and  $I_2$ , generated from the same layout. We adopt the LPIPS metric [48] in computing the DS,

$$DS(I_1, I_2) = \sum_{i=1}^n \frac{1}{|\Lambda_i|} \sum_{\mathbf{p} \in \Lambda_i} \|\omega_i \odot (x_1^i(\mathbf{p}) - x_2^i(\mathbf{p}))\|_2^2, \quad (17)$$

where  $n$  layers of unit-normalized features (in channel dimension),  $x^i$ 's, from a pre-trained DNN (e.g., the VGG network [54]) are used,  $|\Lambda_i|$  the spatial area of a feature map, and  $\omega_i$  the learned parameters by the LPIPS method. So, the higher the DS is, the better a generator model is. Similarly, the mean  $\pm$  std evaluation across multiple runs is used.

The *Classification Accuracy Score* (CAS) [49]. One long-term goal of generative learning in practice is to leverage synthesized images in training discriminative models. The CAS aims to verify how well a classification model trained only on synthesized images can perform on real testing images. So, the higher the CAS is, the better a generator model is. In contrast to the CAS, the classification accuracy metric used in the Layout2Im [20] is based on models trained with real image and tested on synthesized images, which may overlook the diversity of synthesized images.

### 3.4 Quantitative results

We first present the quantitative results and analyses. Table 1 and Table 2 summarize comparisons in terms of the our metrics in the two datasets.

At the resolution of  $64 \times 64$ , our LostGAN-V1 obtains the best performance in comparison. It obtains slightly better Inception Score in both datasets and FID in the VG dataset than the Layout2Im. It obtains significantly better FID in the COCO-Stuff dataset (by more than 5 points reduction) and DS in both datasets. The diversity score of our LostGAN-V1 outperforms the Layout2Im by relative 288.9% and 277.8% in the two datasets respectively. There are a few other methods tested at the resolution of  $64 \times 64$  in the Layout2Im [20], including the pix2pixHD [14], BicycleGAN [58] and GauGAN [15], which are outperformed by the Layout2Im and thus not included here for the clarity of the table.

At the resolution of  $128 \times 128$ , our LostGAN-V1 obtains better results than the Grid2Im method in the COCO-Stuff dataset, especially by more than 33 points reduction in FID and by relative 42.9% increase in DS. Our LostGAN-V2 further improves the results of LostGAN-V1, except for the IS and DS in the VG dataset. **Remarks.** We empirically observed that the VG dataset includes more diverse object

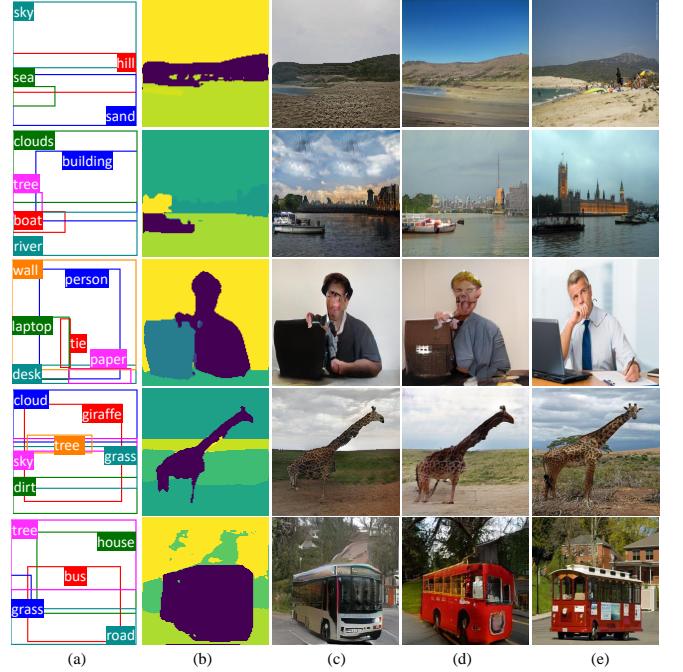


Fig. 8. Comparison of our method and GauGAN [15] at the resolution of  $256 \times 256$ . (a) Input Layouts, (b) Masks learned by our model, (c) Synthesized images by GauGAN using the masks in (b), (d) Generated images by our LostGAN-V2, (e) Ground Truth images. Our model achieve comparable visual performance with GauGAN, which is trained with supervision of ground truth masks. See text for details.

configurations (e.g., bounding boxes may severely overlap in an image such as those for people, cloth and pants), and in general the bounding box annotations in the VG dataset are of lower quality than those in the COCO-Stuff dataset (e.g., they may have significant offsets for certain object instances). Those factors may affect the layout-to-mask component, especially the module of predicting masks from feature maps in the generator, which we think is the reason of LostGAN-V1 slightly outperforming LostGAN-V2 in the VG dataset. Similarly, Layout2Im+OWA [20] suffers a slight drop of performance in the VG dataset after introducing an object-wise attention mechanism to model shape of different objects. Considering those, we only test our LostGAN-V2 at higher resolutions than  $128 \times 128$ .

At the resolution of  $256 \times 256$ , our LostGAN-V2 also obtains better results than the Grid2Im method by more than 2 points increase in IS, 23 points reduction in FID and relative 61.8% increase in DS in the COCO-Stuff dataset.

At the resolution of  $512 \times 512$ , there is no results from other baselines. Our LostGAN-V2 obtains better DS than the DS at the resolution of  $256 \times 256$ . However, our LostGAN-V2 obtains slightly worse results than those obtained at the resolution of  $256 \times 256$  in terms of IS and FID. This phenomenon has been also observed in the BigGAN [5], which indicates, on the one hand, that more research are entailed to improve the quality of high resolution image synthesis, and on the other hand, that the models (Inception V3 pretrained in ImageNet at the resolution of  $300 \times 300$ ) used in computing IS and FID may need to change.

To compare the CAS, we train the ResNet-101 [24] on cropped and resized objects at a resolution of  $32 \times 32$  from generated images (five samples generated for each

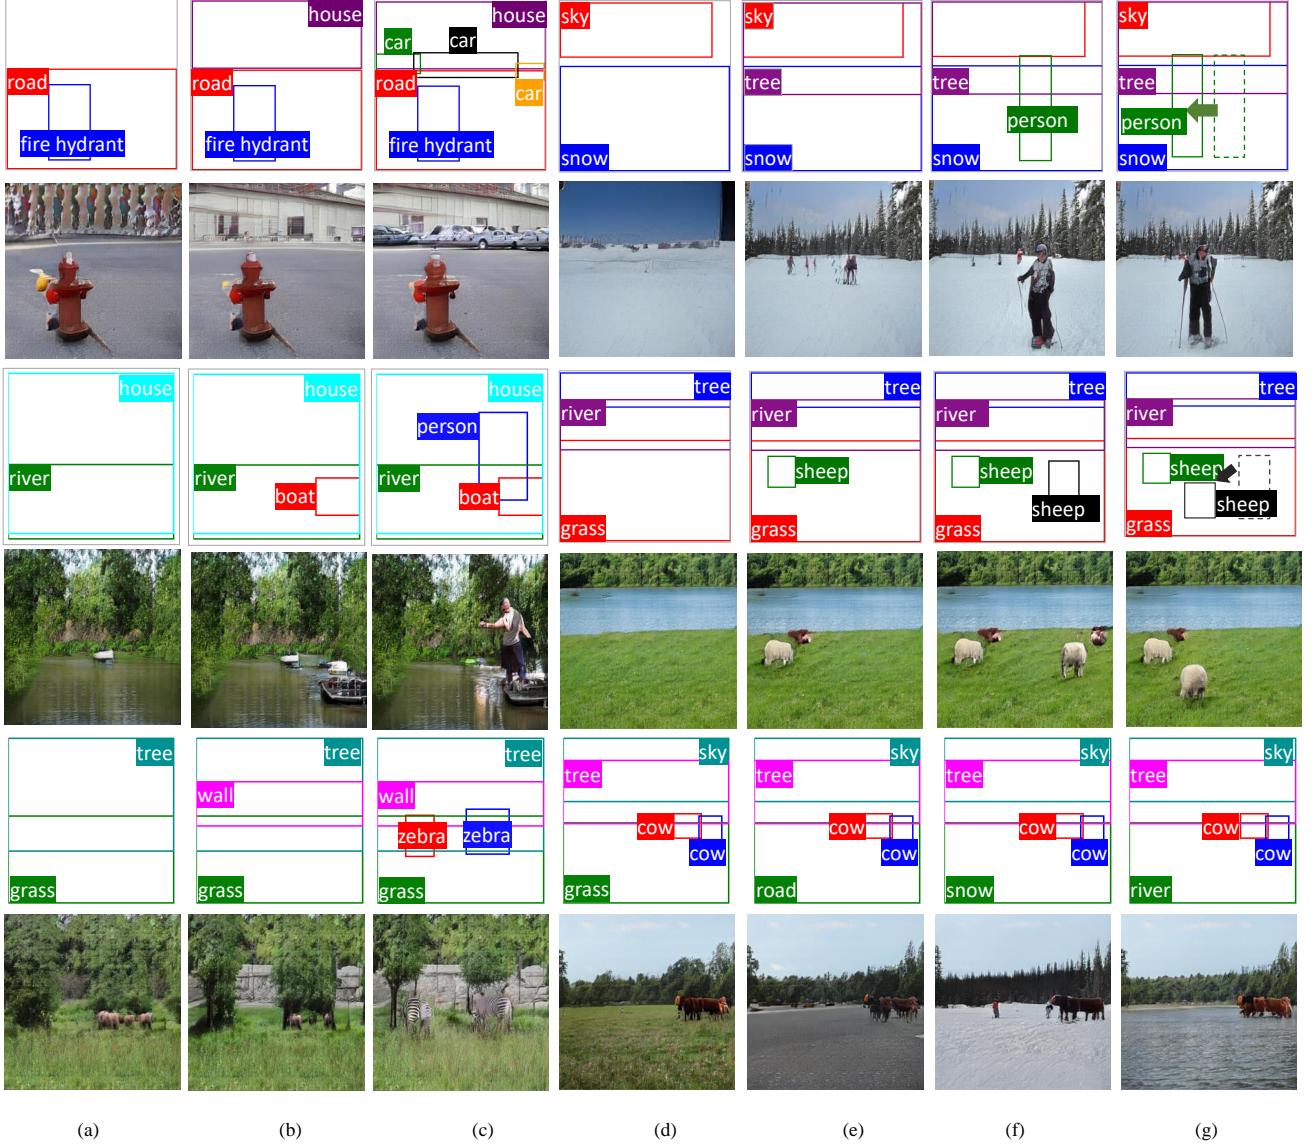


Fig. 9. Layout Control in our LostGAN-V2: image synthesis results by adding new objects, changing the spatial position or the category label of a bounding box in a layout. Best viewed in magnification and color.

layout in the testing set) and evaluate the trained model on objects cropped and resized from real testing images. We follow the widely used settings of ResNet-101 on the CIFAR-10/100 (with images at the resolution  $32 \times 32$ ). We train a 171-category classification ResNet-101 in the COCO-Stuff dataset and a 178-category ResNet-101 in the VG dataset. For synthesized images at the three resolutions, our LostGANs obtain the best accuracy, often by large margin. These results are aligned with the higher DS results consistently obtained by our methods. Hopefully, with more research in the future work, we will be able to generate high-fidelity and high-resolution images from reconfigurable layouts and styles to facilitate more powerful discriminative learning, especially for handling some long-tail or corner situations.

We also compare with the state-of-the-art semantic-map-to-image method, the GauGAN [15]. Instead of using ground-truth semantic maps, we use the masks learned by our LostGAN-V2. Fig. 8 shows some examples, from which we can see the generator in our LostGAN-V2 works rea-

TABLE 3  
Quantitative comparisons using the Inception Score (IS, higher is better), the FID (lower is better) and Diversity Score (DS, higher is better) evaluation on COCO-Stuff dataset at the resolution of  $256 \times 256$ . See text for details.

| Methods                | IS↑              | FID↓  | DS↑             |
|------------------------|------------------|-------|-----------------|
| Our Mask + GauGAN [15] | $19.35 \pm 0.73$ | 41.11 | $0.38 \pm 0.12$ |
| LostGAN-V2             | $18.01 \pm 0.5$  | 42.55 | $0.55 \pm 0.09$ |

sonably good, comparing to the GauGAN that are trained with ground-truth masks. Table 3 shows the comparisons in terms of IS, FID and DS. GauGAN obtains slightly better IS and FID than our LostGAN-V2, while our LostGAN-V2 achieves better DS.

### 3.5 Qualitative results

Fig. 6 and 7 show images synthesized by different models from the same layout in COCO-Stuff and VG respectively.

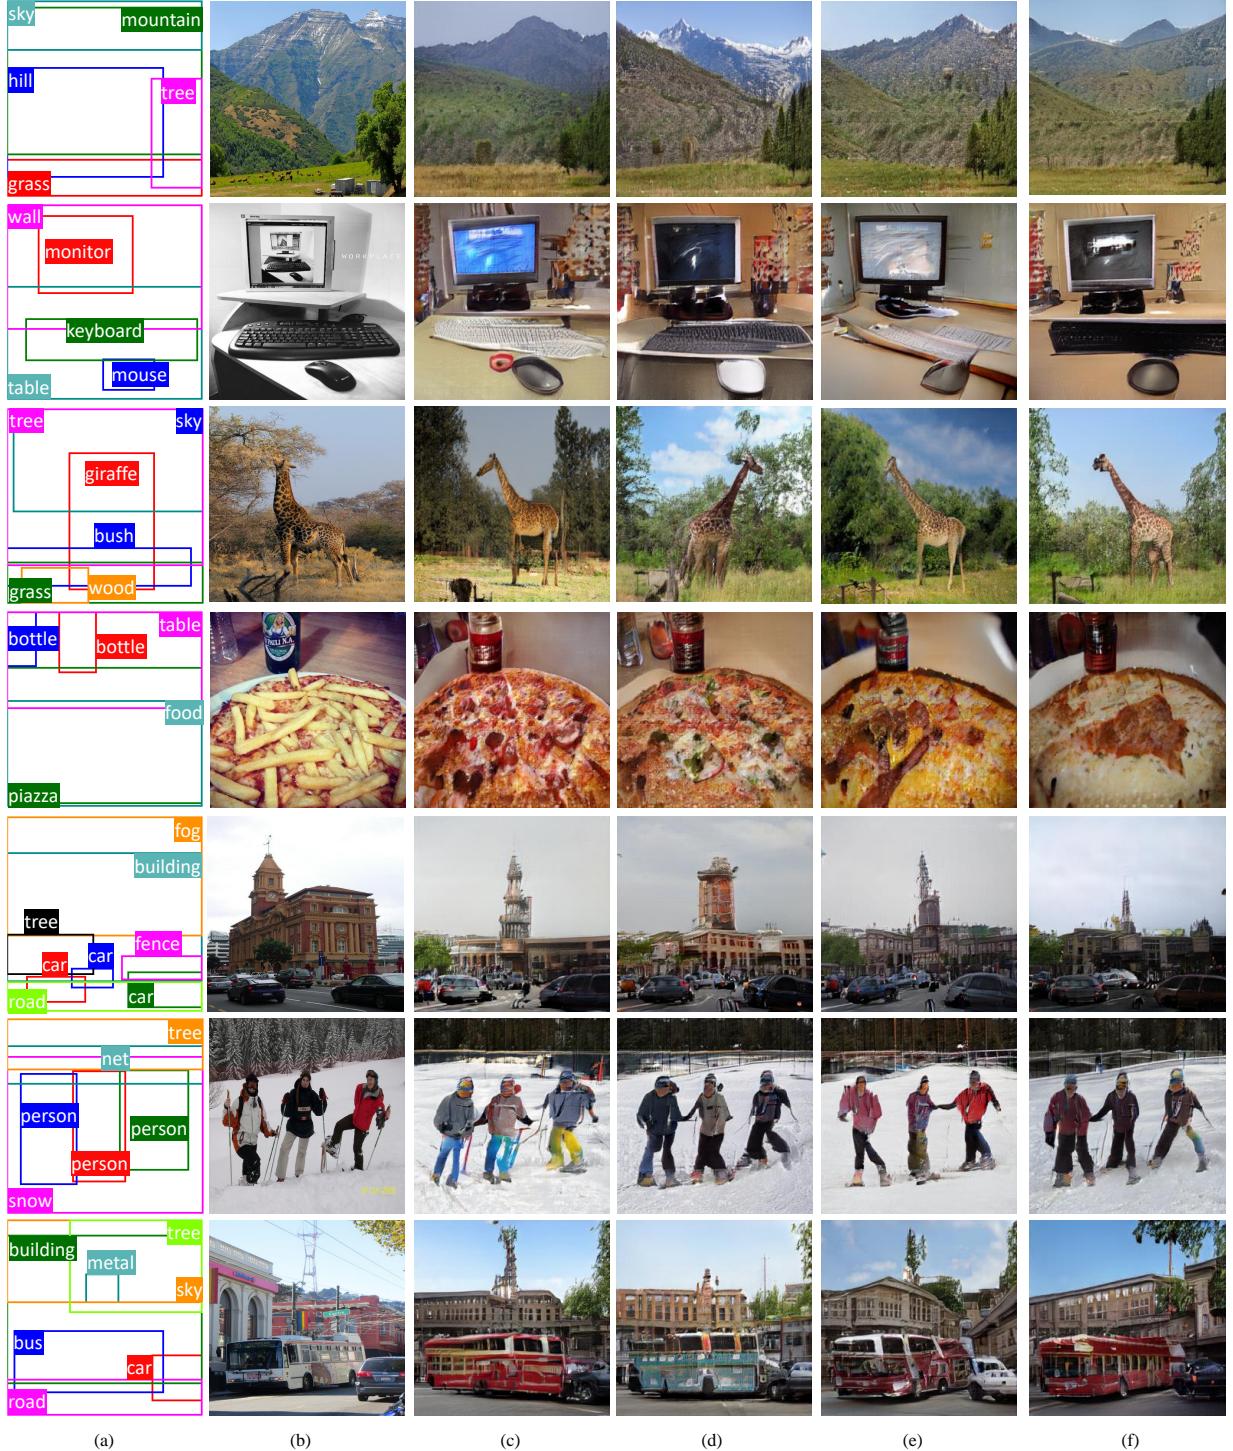


Fig. 10. Style Control: multiple samples generated from the same layout with different styles. Synthesized images have various visual appearance while preserving objects at desired locations. (a) Layout, (b) GT Image and (c-f) Synthesized images by our LostGAN-V2 256 × 256.

The input layouts are quite complex. Our LostGAN-V2 can generate visually more appealing images with more recognizable objects that are consistent with input layouts at resolution 256×256. We show more examples of layout and style control in our LostGAN-V2, in addition to Fig. 2.

In Fig. 9, **layout control** is demonstrated by adding object to, or moving a bounding box in a layout. When adding extra objects or moving the bounding box of one instance, our model can generate reasonable objects at desired

position while keeping existing objects unchanged as we keep the input style of existing objects fixed. When moving the bounding box of one object, style of generated object in new position can also be kept consistent, e.g., in the top-right of Fig. 9, the person is moved while keep style feature like pose and color of clothes unaffected.

In Fig. 10, we show **image-level style control** of our model by synthesizing images with different visual appearance for a given layout while preserving objects at desired



Fig. 11. Linear interpolation of object instance style in our LostGAN-V2. We first generate (b) and (g) using different object style latent codes, and then synthesize (c)-(f) using the style code linearly interpolated from those in (b) and (f). From top to bottom: interpolation of style in grass, sky, cow, bus and giraffe.

locations. In Fig. 11, we show **object-level style control** of our model by gradually morphing styles of one instance in different images.

In Fig. 12, we show some selected examples of synthesized images at the resolution of  $512 \times 512$ . We observed that it is more difficult to generate realistic looking images at the resolution of  $512 \times 512$ .

### 3.6 Ablation Study

We conduct an ablation study on the mask generation component. Fig. 13 shows examples of learned masks, in which even for complex scene with multiple overlapping objects, synthesized images and learned masks are consistent and semantically reasonable. Compared to the input bounding boxes, the learned masks help reduce the semantic gap in layout-to-image. Those masks are learned jointly with image synthesis in a single generator in a weakly-supervised manner, verifying our proposed pipeline of simultaneously learning layout-to-mask-to-image.

Fig. 14 shows examples of mask refinement in the process of generation. The initial mask generation can produce reasonably good results, which are refined in the cascade of

integrating masks learned from feature maps, especially for object boundaries (e.g., comparing (b) and (f)). This mask refinement is the main technical improvement between our LostGAN-V1 and LostGAN-V2, which also verifies the overall improvement by the updated LostGAN-V2 in experiments. To further investigate the effects of mask refinement, after training, we compare the performance of different models with some of mask refinement stages removed. As shown in Table 4, the last row shows the full model with all the mask components,  $m_0 \dots m_5$ . In a backward way, if we remove the mask refinement stage by stage in the generator, the performance (Inception Score and FID) are indeed negatively affected. However, if we remove all the mask refinement stages and only use the initial masks, the performance is better than the model with mask refinement in the first stage,  $m_0 + m_1$ . One potential reason is that the resolution of first stage is very low, from which the learned masks may overlook objects of small sizes and introduce artifacts in the predicted masks. After observing this in the ablation study, we re-trained a model without using  $m_1$  in COCO-Stuff and did not observe performance improvement, so we did not re-train all the models used in



Fig. 12. Some selected examples of synthesized images at the resolution of  $512 \times 512$  in COCO-Stuff by our LostGAN-V2.

TABLE 4

Effects of the mask refinement in our LostGAN-V2  $256 \times 256$  in COCO-Stuff.  $m_0$  represents initial masks generated from the joint label and style encoding.  $m_i$  represents the refined masks at the  $i$ -th stage of the generator. See text for details.

| Mask branch     | IS $\uparrow$    | FID $\downarrow$ |
|-----------------|------------------|------------------|
| $m_0$           | $16.68 \pm 0.42$ | 48.54            |
| $m_0 + m_1$     | $14.14 \pm 0.33$ | 63.96            |
| $m_0 \dots m_2$ | $17.10 \pm 0.56$ | 48.94            |
| $m_0 \dots m_3$ | $17.46 \pm 0.34$ | 44.38            |
| $m_0 \dots m_4$ | $17.51 \pm 0.41$ | 42.49            |
| $m_0 \dots m_5$ | $18.01 \pm 0.50$ | 42.55            |

our experiments.

TABLE 5

mIoU between masks and their nearest neighbor in ground truth masks.

| person | car     | plane | bus      | train  | truck | boat        |
|--------|---------|-------|----------|--------|-------|-------------|
| 53.8   | 66.5    | 58.0  | 75.0     | 70.8   | 66.1  | 63.1        |
| zebra  | hydrant | pizza | elephant | laptop | bench | <b>mean</b> |
| 66.9   | 59.2    | 77.7  | 62.3     | 57.0   | 62.8  | 56.5        |

To investigate the quality of learned masks, we resort to the intersection-over-union (IoU) metric used in object semantic segmentation. We measure the IoU performance in the COCO-Stuff training dataset. We first crop masks for each category and then resize all the masks to the same resolution of  $32 \times 32$ . After training the LostGAN-V2  $256 \times 256$ , we run inference on each layout in the training dataset (one

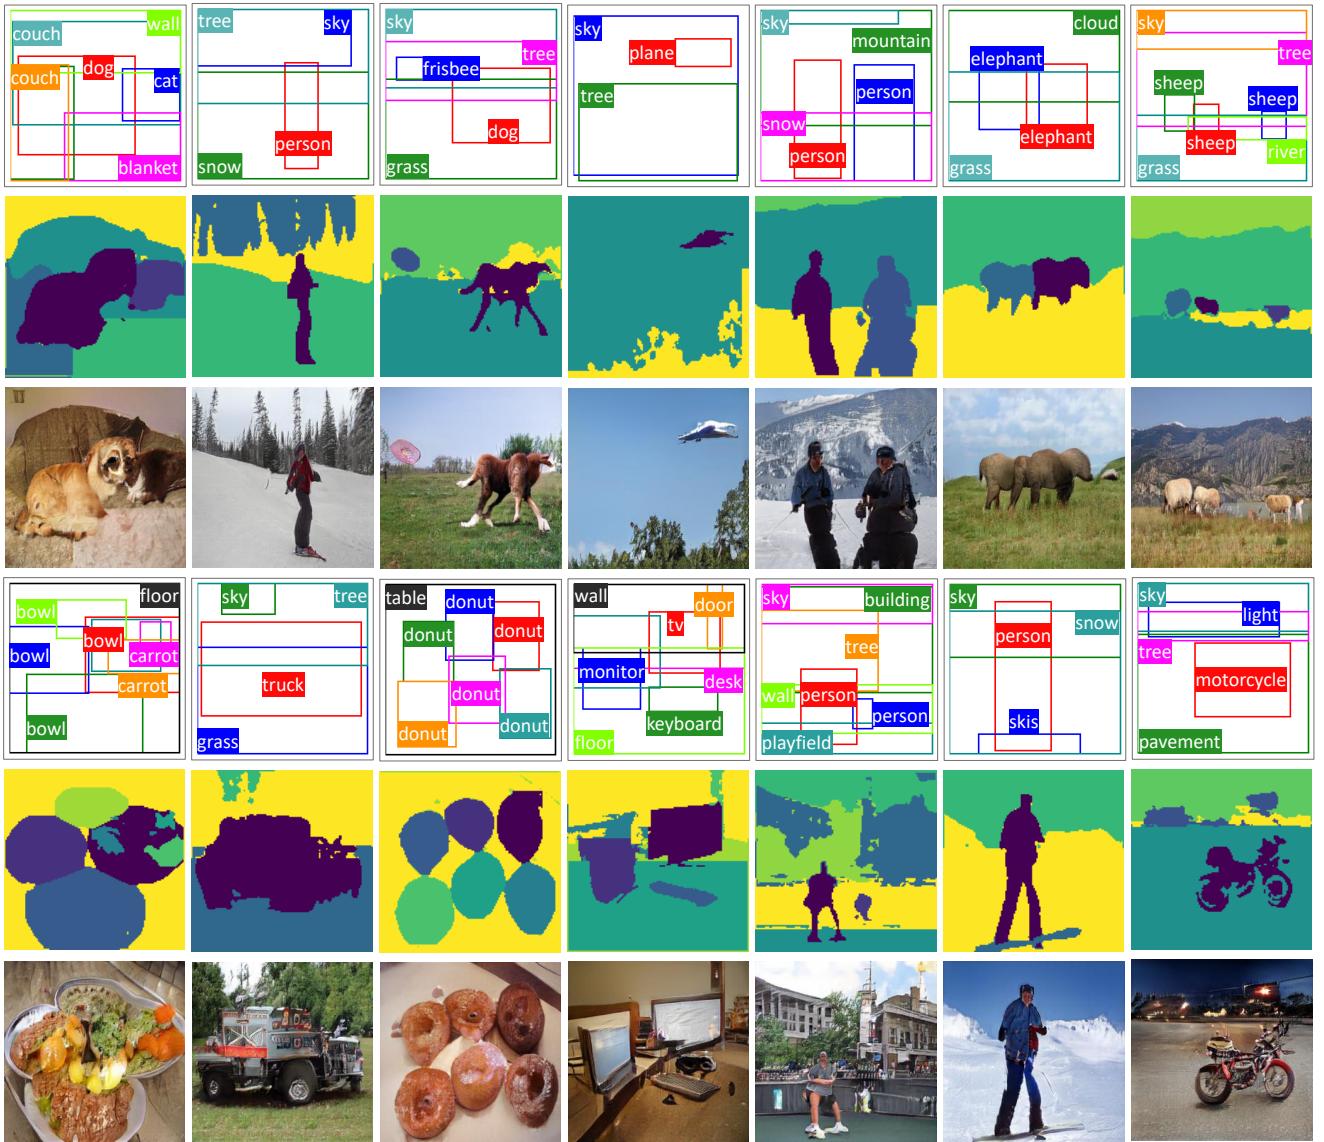


Fig. 13. Examples of synthesized images and learned masks by our LostGAN-V2  $256 \times 256$ .

run is used for simplicity) and obtain the learned masks. We then crop and resize object masks in the same way as done for the ground-truth object masks. For each learned object mask, we retrieve the top- $k$  nearest neighbors in terms of mask IoU in the set of ground-truth object masks. Fig. 15 shows four examples with top-10 nearest neighbors. Table 5 shows the mean IoUs for 13 selected object categories which have reasonably high IoUs.

### 3.7 Failure Examples

Fig. 16 shows some failure examples. We observe that our proposed model is not capable of capturing interactions between person and small objects, e.g., person and baseball bat, tennis racket, handbag, etc. From the learned masks, we can also see why the model can not synthesize good images. We leave this to our future work by investigating methods of learning fine-grained part-level masks.

*Remarks.* The generative learning of layout-to-image is still at a early stage of development in terms of synthesizing

high-fidelity images, compared to the results of BigGANs [5] in ImageNet and StyleGANs [8] for faces. Overall, we can observe the quality of image generation from layout is still not sufficiently good, especially for articulated objects (such as people) and fine-grained object-object interactions at high resolution (*e.g.*, examples in Fig. 16). In the meanwhile, we also note that the differences between the goals of BigGANs and StyleGANs and those of controllable layout-to-image are non-trivial. For example, we can use a trained BigGAN to generate cat images, and as long as the generated images look realistic and sharp with one or more than one cats, we shall think it does a great job (without requiring how many cats should appear and where they should be). Similarly, we can train a StyleGAN to generate face images, and we shall be happy if realistic and sharp face images are generated with a natural style (*e.g.*, smiling or sad). Controllable layout-to-image has more fine-grained requirements. Those being said, based on the promising results of GauGANs [15] using annotated semantic maps in image synthesis, we think

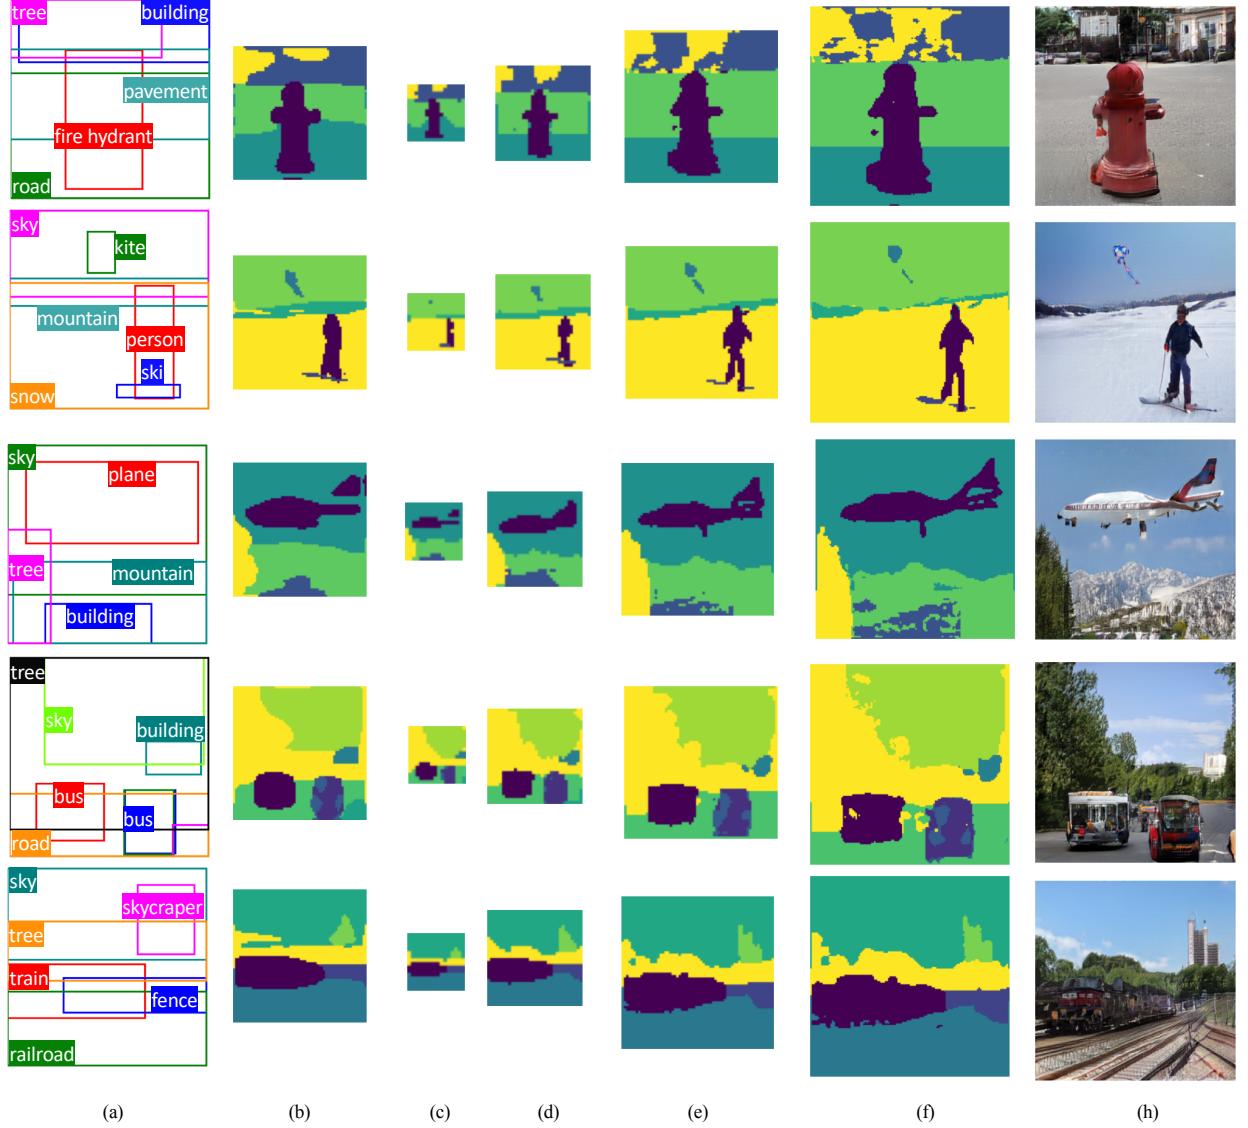


Fig. 14. Examples of mask refinement in the generator. (a) Layouts, (b) Initial masks generated from the joint label and style encoding, (c-f) Mask refinement using masks generated from feature maps at different stages in the generator, (h) Synthesized images.

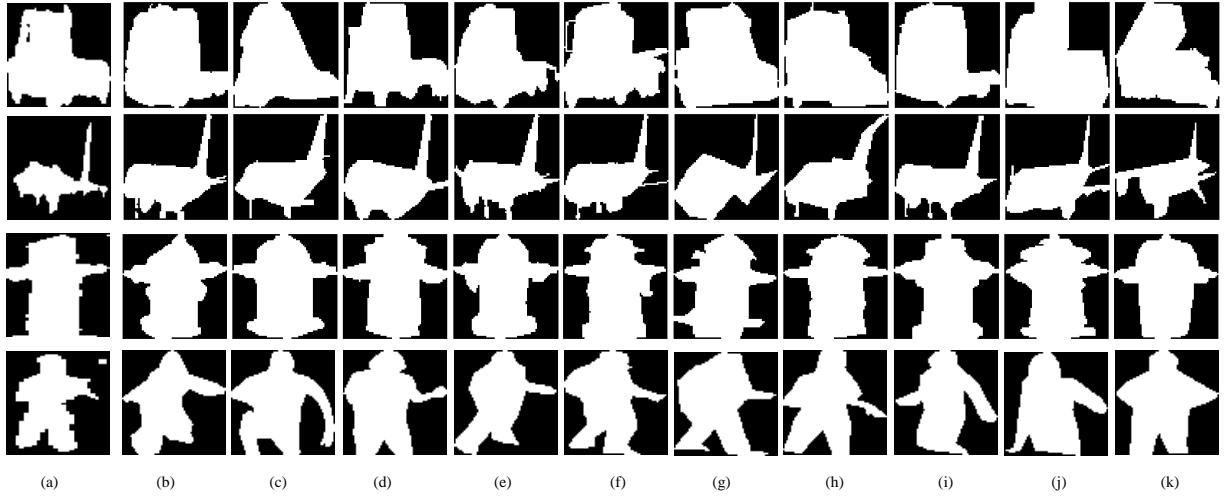


Fig. 15. Examples of learned masks and their nearest neighbors in the ground-truth masks in the COCO-Stuff training dataset: truck, airplane, hydrant and person (from top to bottom). (a) Masks learned by our LostGAN-V2 256 × 256 and (b-k) top-10 nearest neighbors. All masks are cropped and resized to the resolution of 32 × 32. See text for details.

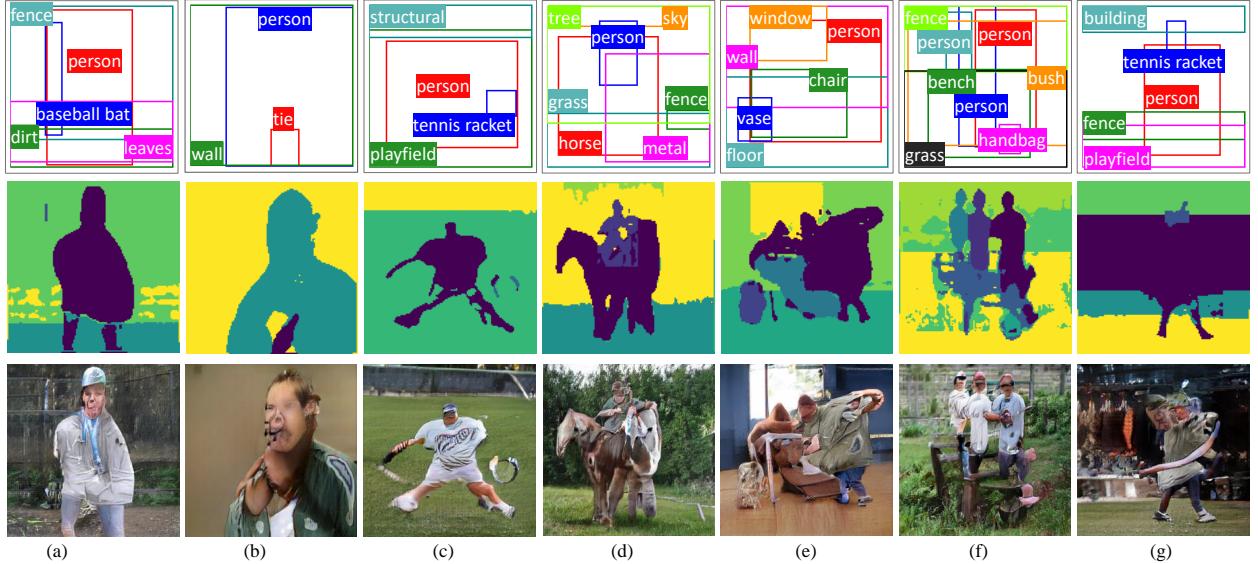


Fig. 16. Examples of failure cases of our LostGAN-V2 256 × 256.

the proposed layout-to-mask-to-image pipeline and LostGAN worth further explorations of seeking more powerful weakly-supervised learning of layout-to-mask. For example, we can develop more sophisticated mask generators and the “ToMask” modules in Fig. 4, and explore different consistency constraints between the “ToMask” modules, similar to the recently proposed PointRend method for improving Mask-RCNNs [59]. We leave those for the future work.

## 4 CONCLUSIONS

This paper studies the generative learning problem of layout-to-image with a focus on controllable image synthesis from reconfigurable layouts and styles. We first propose an intuitive pipeline of learning layout-to-mask-to-image. We then present a layout- and style-based architecture for generative adversarial networks (termed LostGANs). Our proposed LostGAN can be trained end-to-end to generate images from reconfigurable layout and style with strong style and layout controllability. Our proposed LostGAN can also learn fine-grained object masks in a weakly-supervised manner to bridge the gap between layouts and images by a novel object instance-sensitive layout-aware feature normalization (ISLA-Norm) scheme. State-of-the-art performance is obtained in the COCO-Stuff and Visual Genome datasets.

## ACKNOWLEDGEMENTS

This work was supported in part by NSF IIS-1909644 and ARO Grant W911NF1810295. The views presented in this paper are those of the authors and should not be interpreted as representing any funding agencies.

## REFERENCES

- [1] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, “Generative adversarial nets,” in *Advances in neural information processing systems*, 2014, pp. 2672–2680.
- [2] A. Radford, L. Metz, and S. Chintala, “Unsupervised representation learning with deep convolutional generative adversarial networks,” in *International Conference on Learning Representations (ICLR)*, 2016.
- [3] H. Zhang, I. Goodfellow, D. Metaxas, and A. Odena, “Self-attention generative adversarial networks,” in *Proceedings of the 36th International Conference on Machine Learning*, 2019, pp. 7354–7363.
- [4] T. Miyato, T. Kataoka, M. Koyama, and Y. Yoshida, “Spectral normalization for generative adversarial networks,” in *International Conference on Learning Representations*, 2018.
- [5] A. Brock, J. Donahue, and K. Simonyan, “Large scale GAN training for high fidelity natural image synthesis,” in *International Conference on Learning Representations*, 2019.
- [6] T. Miyato and M. Koyama, “cGANs with projection discriminator,” in *International Conference on Learning Representations*, 2018.
- [7] T. Karras, T. Aila, S. Laine, and J. Lehtinen, “Progressive growing of GANs for improved quality, stability, and variation,” in *International Conference on Learning Representations*, 2018.
- [8] T. Karras, S. Laine, and T. Aila, “A style-based generator architecture for generative adversarial networks,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2019, pp. 4401–4410.
- [9] A. Odena, C. Olah, and J. Shlens, “Conditional image synthesis with auxiliary classifier gans,” in *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, 2017, pp. 2642–2651.
- [10] T. Kim, M. Cha, H. Kim, J. K. Lee, and J. Kim, “Learning to discover cross-domain relations with generative adversarial networks,” in *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, 2017, pp. 1857–1865.
- [11] P. Isola, J.-Y. Zhu, T. Zhou, and A. A. Efros, “Image-to-image translation with conditional adversarial networks,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2017, pp. 1125–1134.
- [12] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros, “Unpaired image-to-image translation using cycle-consistent adversarial networks,” in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 2223–2232.
- [13] X. Huang, M.-Y. Liu, S. Belongie, and J. Kautz, “Multimodal unsupervised image-to-image translation,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, 2018, pp. 172–189.
- [14] T.-C. Wang, M.-Y. Liu, J.-Y. Zhu, A. Tao, J. Kautz, and B. Catanzaro, “High-resolution image synthesis and semantic manipulation with conditional gans,” in *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2018.
- [15] T. Park, M.-Y. Liu, T.-C. Wang, and J.-Y. Zhu, “Semantic image synthesis with spatially-adaptive normalization,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2019, pp. 2337–2346.

- [16] S. Reed, Z. Akata, X. Yan, L. Logeswaran, B. Schiele, and H. Lee, "Generative adversarial text-to-image synthesis," in *Proceedings of The 33rd International Conference on Machine Learning*, 2016.
- [17] H. Zhang, T. Xu, H. Li, S. Zhang, X. Wang, X. Huang, and D. N. Metaxas, "Stackgan: Text to photo-realistic image synthesis with stacked generative adversarial networks," in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 5907–5915.
- [18] J. Johnson, A. Gupta, and L. Fei-Fei, "Image generation from scene graphs," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 1219–1228.
- [19] O. Ashual and L. Wolf, "Specifying object attributes and relations in interactive scene generation," in *Proceedings of the IEEE International Conference on Computer Vision*, 2019, pp. 4561–4569.
- [20] B. Zhao, W. Yin, L. Meng, and L. Sigal, "Layout2image: Image generation from layout," *International Journal of Computer Vision*, pp. 1–18, 2020.
- [21] T. Xu, P. Zhang, Q. Huang, H. Zhang, Z. Gan, X. Huang, and X. He, "AttnGAN: Fine-grained text to image generation with attentional generative adversarial networks," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 1316–1324.
- [22] H. Caesar, J. Uijlings, and V. Ferrari, "Coco-stuff: Thing and stuff classes in context," in *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2018.
- [23] S. Hong, D. Yang, J. Choi, and H. Lee, "Inferring semantic layout for hierarchical text-to-image synthesis," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 7986–7994.
- [24] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 770–778.
- [25] K. He, G. Gkioxari, P. Dollár, and R. Girshick, "Mask r-cnn," in *Proceedings of the IEEE international conference on computer vision*, 2017, pp. 2961–2969.
- [26] W. Sun and T. Wu, "Image synthesis from reconfigurable layout and style," in *Proceedings of the IEEE International Conference on Computer Vision*, 2019, pp. 10531–10540.
- [27] D. Tran, R. Ranganath, and D. Blei, "Hierarchical implicit models and likelihood-free variational inference," in *Advances in Neural Information Processing Systems 30*, 2017, pp. 5523–5533.
- [28] J. H. Lim and J. C. Ye, "Geometric gan," *arXiv preprint arXiv:1705.02894*, 2017.
- [29] S. Ioffe and C. Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift," in *Proceedings of the 32nd International Conference on Machine Learning*, 2015, pp. 448–456.
- [30] A. V. Oord, N. Kalchbrenner, and K. Kavukcuoglu, "Pixel recurrent neural networks," in *Proceedings of The 33rd International Conference on Machine Learning*, 2016, pp. 1747–1756.
- [31] A. Van den Oord, N. Kalchbrenner, L. Espeholt, O. Vinyals, A. Graves et al., "Conditional image generation with pixelcnn decoders," in *Advances in neural information processing systems*, 2016, pp. 4790–4798.
- [32] D. P. Kingma and M. Welling, "Auto-encoding variational bayes," in *Proceedings of the 2nd International Conference on Learning Representations (ICLR)*, 2014.
- [33] Y. Pu, Z. Gan, R. Henao, X. Yuan, C. Li, A. Stevens, and L. Carin, "Variational autoencoder for deep learning of images, labels and captions," in *Advances in neural information processing systems*, 2016, pp. 2352–2360.
- [34] M. Arjovsky, S. Chintala, and L. Bottou, "Wasserstein generative adversarial networks," in *Proceedings of the 34th International Conference on Machine Learning*, 2017, pp. 214–223.
- [35] T. Han, E. Nijkamp, X. Fang, M. Hill, S. Zhu, and Y. N. Wu, "Divergence triangle for joint training of generator model, energy-based model, and inferential model," in *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2019, Long Beach, CA, USA, June 16–20, 2019*, pp. 8670–8679.
- [36] M. Mirza and S. Osindero, "Conditional generative adversarial nets," *arXiv preprint arXiv:1411.1784*, 2014.
- [37] M.-Y. Liu, X. Huang, A. Mallya, T. Karras, T. Aila, J. Lehtinen, and J. Kautz, "Few-shot unsupervised image-to-image translation," in *Proceedings of the IEEE International Conference on Computer Vision*, 2019, pp. 10551–10560.
- [38] E. Mansimov, E. Parisotto, J. L. Ba, and R. Salakhutdinov, "Generating images from captions with attention," in *International Conference on Learning Representations (ICLR)*, 2016.
- [39] W. Li, P. Zhang, L. Zhang, Q. Huang, X. He, S. Lyu, and J. Gao, "Object-driven text-to-image synthesis via adversarial training," in *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019.
- [40] L. Yikang, T. Ma, Y. Bai, N. Duan, S. Wei, and X. Wang, "Pastegan: A semi-parametric method to generate image from scene graph," in *Advances in Neural Information Processing Systems*, 2019, pp. 3950–3960.
- [41] V. Dumoulin, I. Belghazi, B. Poole, O. Mastropietro, A. Lamb, M. Arjovsky, and A. Courville, "Adversarially learned inference," in *International Conference on Learning Representations(ICLR)*, 2017.
- [42] H. De Vries, F. Strub, J. Mary, H. Larochelle, O. Pietquin, and A. C. Courville, "Modulating early visual processing by language," in *Advances in Neural Information Processing Systems*, 2017, pp. 6594–6604.
- [43] V. Dumoulin, J. Shlens, and M. Kudlur, "A learned representation for artistic style," *International Conference on Learning Representations (ICLR)*, vol. 2, 2017.
- [44] T. Hinz, S. Heinrich, and S. Wermter, "Generating multiple objects at spatially distinct locations," in *International Conference on Learning Representations*, 2019.
- [45] ———, "Semantic object accuracy for generative text-to-image synthesis," *arXiv preprint arXiv:1910.13321*, 2019.
- [46] T. Salimans, I. Goodfellow, W. Zaremba, V. Cheung, A. Radford, and X. Chen, "Improved techniques for training gans," in *Advances in neural information processing systems*, 2016, pp. 2234–2242.
- [47] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, and S. Hochreiter, "Gans trained by a two time-scale update rule converge to a local nash equilibrium," in *Advances in Neural Information Processing Systems*, 2017, pp. 6626–6637.
- [48] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, and O. Wang, "The unreasonable effectiveness of deep features as a perceptual metric," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 586–595.
- [49] S. Ravuri and O. Vinyals, "Classification accuracy score for conditional generative models," in *Advances in Neural Information Processing Systems*, 2019, pp. 12247–12258.
- [50] R. Krishna, Y. Zhu, O. Groth, J. Johnson, K. Hata, J. Kravitz, S. Chen, Y. Kalantidis, L.-J. Li, D. A. Shamma et al., "Visual genome: Connecting language and vision using crowdsourced dense image annotations," *International Journal of Computer Vision*, vol. 123, no. 1, pp. 32–73, 2017.
- [51] T.-Y. Lin, P. Dollar, R. Girshick, K. He, B. Hariharan, and S. Belongie, "Feature pyramid networks for object detection," in *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, July 2017.
- [52] S. Ren, K. He, R. Girshick, and J. Sun, "Faster r-cnn: Towards real-time object detection with region proposal networks," in *Advances in Neural Information Processing Systems 28*, 2015, pp. 91–99.
- [53] J. Johnson, A. Alahi, and L. Fei-Fei, "Perceptual losses for real-time style transfer and super-resolution," in *European conference on computer vision*. Springer, 2016, pp. 694–711.
- [54] K. Simonyan and A. Zisserman, "Very deep convolutional networks for large-scale image recognition," in *International Conference on Learning Representations*, 2015.
- [55] A. M. Saxe, J. L. McClelland, and S. Ganguli, "Exact solutions to the nonlinear dynamics of learning in deep linear neural networks," in *International Conference on Learning Representations(ICLR)*, 2014.
- [56] D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," in *International Conference on Learning Representations(ICLR)*, 2015.
- [57] D. Dowson and B. Landau, "The fréchet distance between multivariate normal distributions," *Journal of multivariate analysis*, vol. 12, no. 3, pp. 450–455, 1982.
- [58] J.-Y. Zhu, R. Zhang, D. Pathak, T. Darrell, A. A. Efros, O. Wang, and E. Shechtman, "Toward multimodal image-to-image translation," in *Advances in Neural Information Processing Systems 30*, 2017, pp. 465–476.
- [59] A. Kirillov, Y. Wu, K. He, and R. B. Girshick, "Pointrend: Image segmentation as rendering," in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.